
Detecting AI Training in Datacentre Power

Tom Kimpson¹ Mauricio Baker² Emlyn Graham³ Anjay Friedman² Coby Joseph²

Abstract

Compute-governance regimes that restrict large training runs presuppose that a regulator can tell when training happens on monitored hardware, and a power meter outside the prover’s control is the one passive channel available against a prover who does not cooperate. We ask what such a meter can certify, and answer with a ladder of four claims, each with a stated information cost and a hard ceiling. Evidence comes from a literature-parameterized scenario model of datacentre AI workloads and an explicit off-chip observation model, against both a benign prover and an adaptive one who knows the detector. The iteration cadence of synchronous training is recoverable by a tracked cyclostationary detector, but the complete pipeline admits no analytic calibration, so we report it as an effective score with per-trace surrogate calibration rather than a constant false-alarm-rate test. Our central result is the de-periodicisation frontier, which prices hiding against measured throughput anchors and is strongly asymmetric: varying real work per iteration hides from every fixed-frequency test at essentially zero cost, while the best anchored attack against the tracking class buys little hiding for a 159% overhead. Two limits are firmer than the detectors. The observation channel is itself a condition on every claim — a one-second-integrating, one-hertz meter defeats all detectors on an honest workload — and the meter certifies a physical schedule, not training semantics: five non-training loads, including one computing gradients and discarding them, are classified as training by every rule we test. Escaping that ceiling means leaving the passive channel for an active challenge protocol, which we specify but do not run.

1. Introduction

AI governance agreements will require regulating training runs (Shavit, 2023; Sastry et al., 2024; Baker et al., 2025; Heim & Koessler, 2024). This will require an external verifier to identify when training is occurring on a monitored cluster, against an operator who is trying to hide their training. There are multiple channels that could provide the verifier with diagnostic information to identify hidden training. These include on-chip mechanisms that log activities for later analysis, network taps that monitor data exchanged between chips, or human-centred approaches such as whistleblower programs or interviews of operator personnel. However each of these channels requires trusting either hardware that is under the control of the operator, or the operator personnel themselves. Conversely, off-chip sensors that take analog, physical measurements of the cluster (e.g. power or electromagnetic emissions) do not require trusting the operator. Such sensors could be retrofitted to existing datacentres (e.g. clamping a power meter on an input cable) and secured against tampering by the operator.

AI training has already been shown to leave a signature in the power drawn by the hardware that can serve as a diagnostic to identify and separate training from inference. A 10 MHz current probe clamped around the GPU’s PCIe power supply cables can separate training from inference and from non-AI work at 97% accuracy (Gargiulo & Kulp, 2026), a result robust against many but not all adversarial strategies. On-chip 1 Hz telemetry counters that measure power (among other telemetry: utilisation, memory, etc.) can classify training at 95% accuracy, again a result robust against many but not all adversarial strategies (Rahman & Tajdari, 2026). In practice, any realistic AI governance agreements that use off-chip sensors will rarely have such fine-grained access to the GPU hardware. Power meters that can be retrofitted to existing facilities and secured against tampering are likely to sit at the level of the rack or facility input, with sampling frequencies ~ 10 Hz (i.e. three to four orders of magnitude slower than the board probe) and behind a power-delivery path that low-passes and dilutes the signal. At this level training can be identified and it is large enough that datacentre and grid operators treat it as an engineering control problem to be dealt with (Patel et al., 2023; Ko & Zhu, 2025).

¹School of Mathematics and Statistics, The University of Melbourne, Parkville, VIC 3010, Australia ²RAND Corporation ³Affiliation TBD. Correspondence to: Tom Kimpson <tom.kimpson@unimelb.edu.au>.

In this paper we develop a new method for detecting training signatures given noisy facility level power measurements, against an adversarial operator who is trying to hide their training.

We do:

- * specify an off-chip observation model and a literature-parameterised generator of datacentre training and inference power,
- * construct a detector for the cadence that tracks the wandering line, borrowing techniques from signal analysis and GW astrophysics
- * We construct the Pareto frontier and show the trade-off between how much of the cadence an adaptive prover removes and what that removal costs in measured throughput, and derive an identifiability lower bound on the cost of hiding to a stated detectability (Sections 6 and G.1).
- * We derive a minimum meter specification, the sampling rate, integration time and in-band transfer i.e. the specs needed for a meter to work
- * We restrict attention to synchronous, iteration-structured training observed through a time-resolved trace; the integrated-energy channel, asynchronous schedules and the calibration of the generator to distributed hardware are outside scope, and Sections 7 and 7.2 say what that costs.

Unlike a covert transmitter, a training operator cannot freely reshape its signal: the communication phase has a fixed physical duration, and every distortion of the cadence costs throughput. It is this constraint that makes a cost-of-hiding bound possible.

This paper is organised as follows. Section 2 gives the model of datacentre training and inference power that generates every trace below, and Section 3 fixes what the verifier sees through an off-chip meter and what an adaptive prover may do to hide from it. Section 4 develops the detector, a Viterbi line tracker; Section 5 reports its power against the inference null and the minimum meter it needs; and Section 6 prices what it costs an adaptive prover to defeat it. Section 7 states what a detection does and does not certify, holds the model against measured single-GPU hardware, and reads the results back as a ladder of claims a regulator can buy.

2. A model of datacentre training power

In this section we build a model of the power draw from a data centre running AI workloads. This model will serve as a synthetic data generator to test our methods, in lieu of having access to real empirical data. We adopt the stochastic model of AI training power from Ko & Zhu (2025), developed for the analysis of grid stability, modified to add an Ornstein-

Uhlenbeck drift on the training frequency. The reader is referred to Ko & Zhu (2025) for the full training model and its context.

This section is organised as follows. Section 2.1 describes the power draw from a training workload, following Ko & Zhu (2025), and adds our drift of its cadence. Section 2.2 describes the power draw from a fine-tuning workload. Section 2.3 describes the power draw from an inference workload. Section 2.4 combines these workloads into the load of a whole datacentre: one dominant workload, either training or inference, over a background of smaller trainings and fine-tunings. Table 1 collects every parameter value of the model.

2.1. Training

Multi GPU AI training of the kind that typifies frontier models is composed of a two-phase iterative cycle. First there is a power intensive, high compute “up” phase where the forward/backward pass happens. Second there is a low power, low compute, communication “down” phase where the gradient all reduce happens.

Iteration i has period

$$T_i = \frac{1}{f_0 (1 + \xi_i)}, \quad \xi_i \sim \mathcal{N}(0, \sigma_\xi^2), \quad (1)$$

where f_0 is the baseline fluctuation frequency. For a given training run f_0 is fixed. Between training runs f_0 can vary depending on model size, batch size, parallelism layout, interconnect and hardware (Patel et al., 2024a; Jain et al., 2024). To describe the inter-workload variation, we take $f_0 \sim \mathcal{U}(f_0^{\text{lo}}, f_0^{\text{hi}})$. The period splits into an up-phase of duration $\rho_i T_i$ and a down-phase of duration $(1 - \rho_i) T_i$, with the phase ratio $\rho_i \sim \mathcal{U}(\rho^{\text{lo}}, \rho^{\text{hi}})$ redrawn each iteration. Writing t for time within iteration i , $P(t)$ for the power the run draws above the facility’s static overhead (added in Section 3.1), and P_{peak} for its peak,

$$\frac{P(t)}{P_{\text{peak}}} = \begin{cases} 1 + \Delta_i^{\text{up}} + \nu^{\text{up}}(t), & t < \rho_i T_i, \\ 1 - \Delta_i^{\text{dn}} + \nu^{\text{dn}}(t), & t \geq \rho_i T_i, \end{cases} \quad (2)$$

where the per-phase level offsets are $\Delta_i^{\text{up}} \sim \mathcal{N}(0, \sigma_\Delta^2)$ and $\Delta_i^{\text{dn}} \sim \mathcal{N}(\mu_\Delta, \sigma_\Delta^2)$, and $\nu^{\text{up}}, \nu^{\text{dn}}$ are zero-mean intra-phase Gaussian fluctuations with standard deviations $\kappa_\nu s_i^{\text{up}}$ and $\kappa_\nu s_i^{\text{dn}}$. The s_i^{up} and s_i^{dn} are redrawn each phase, uniformly from the ranges $[s^{\text{up,lo}}, s^{\text{up,hi}}]$ and $[s^{\text{dn,lo}}, s^{\text{dn,hi}}]$, and κ_ν is a fixed scale factor. Table 1 lists the values of every parameter in this section.

Equations (1) and (2) are the exact training workload model from Ko & Zhu (2025). We make one modification, adding a slow drift to the cadence f_0 itself. The jitter ξ_i does not produce drift. Because it is redrawn independently each iteration, the cadence averaged over many iterations is still

f_0 . The timing errors accumulate as a random walk in phase, which broadens the spectral line into a Lorentzian of full width $\approx 2\pi\sigma_\xi^2 f_0$ (to leading order in σ_ξ) but leaves its centre at f_0 . Real runs vary differently. Specifically, the iteration time changes in ways that persist. Every training step waits for the slowest GPU, so one slow component slows the whole run. In a production LLM cluster, these slowdowns last for the whole job rather than varying from step to step (Lin et al., 2025). Hardware can also keep running but degrade: CPU contention, overheating GPUs and congested network links cut throughput by tens of percent for periods of tens of seconds to hours (Wu et al., 2024). Over a day, temperature shifts throughput by a further 1–2% (Grattafiori et al., 2024).

We model this by letting the cadence of iteration i be $f_0 g_i$, where g_i is a discrete Ornstein–Uhlenbeck (first-order autoregressive) process that starts at $g_0 = 1$ and reverts to 1:

$$g_i - 1 = (1 - \theta)(g_{i-1} - 1) + \epsilon_i, \quad \epsilon_i \sim \mathcal{N}(0, \sigma_g^2). \quad (3)$$

The period of Equation (1) becomes

$$T_i = \frac{1}{f_0 g_i (1 + \xi_i)}, \quad (4)$$

which reduces to Equation (1) when $g_i \equiv 1$. The stationary variance of Equation (3) is $\sigma_g^2/(2\theta - \theta^2)$, so we set

$$\sigma_g = \frac{\sigma_f}{f_0} \sqrt{2\theta - \theta^2}, \quad (5)$$

which makes the stationary standard deviation of the instantaneous cadence $f_0 g_i$ equal to σ_f , in Hz. The autocorrelation of g_i at lag k is $(1 - \theta)^k$, so the cadence stays correlated over $\sim 1/\theta$ iterations. Over that span it wanders by σ_f , which can exceed both the line width set by ξ_i and the frequency resolution of a trace. We clip g_i from below at $g_{\min}$ so the period stays finite. We hold θ fixed and treat σ_f as a free parameter, sweeping it in Section 5.

2.2. Fine-tuning

Fine-tuning workloads can be represented using the same structure as in Section 2.1, with different parameters (Table 1). We denote the fine-tuning frequency as f_1 . Compared with training, it has a lower cadence band, a longer tail phase (larger ρ_i), a deeper tail-to-idle swing (larger μ_Δ), and slightly smaller level offsets and intra-phase noise. We apply the cadence drift of Equation (3) only to the dominant run; fine-tuning and background training runs keep a fixed cadence ($\sigma_f = 0$).

2.3. Inference

Ko & Zhu (2025) do not present a model for power draw of inference workloads. We now construct one. We suppose three structures on a constant mean serving power

load B . First, token generation (decode). Serving systems use continuous batching, which adds and removes requests from the batch at every step (Yu et al., 2022; Kwon et al., 2023). A per-step cadence exists, but the batch composition reshuffles every step. Unlike training, there is therefore no characteristic frequency: decode power smears across a broad band with no coherent frequency path. We model it as band-limited noise $b(t)$ of unit RMS over $[f_b^{\text{lo}}, f_b^{\text{hi}}]$, scaled by an amplitude A . Second, expert-load imbalance. Under sparse mixture-of-experts routing, tokens are not spread evenly across the experts (Shazeer et al., 2017), so the work per step varies with the mix of requests. We model this as a slow envelope on the decode power: a unit-variance Ornstein–Uhlenbeck process $u(t)$ with correlation time τ_u , entering at depth m . Amplitude modulation spreads the decode energy into sidebands, which makes the null harder, not easier. Third, prompt processing (prefill). Each request begins with a prompt phase that is brief, compute-bound, and draws more power than token generation (Patel et al., 2024a;b). We model these phases as box pulses $s(\cdot)$ of duration d_p and height a , arriving at times τ_k from a Poisson process of rate λ_p . These memoryless arrivals carry no line and simply raise the broadband in-band level. Together,

$$\frac{P^{\text{inf}}(t)}{P_{\text{peak}}} = B + A [1 + m u(t)] b(t) + \sum_k a s(t - \tau_k). \quad (6)$$

The forms of these structures follow how LLM serving systems schedule work; their amplitudes are our choice (Table 1). We set them so that inference is a loud confound rather than a quiet one. As Figure 1 shows, it delivers more in-band power than training but an order of magnitude less spectral contrast.

2.4. Aggregate

A datacentre rarely runs one job at a time. Following Ko & Zhu (2025) we write the load behind a single meter as the aggregate of one dominant training run over a background of smaller trainings and fine-tunings,

$$P_{\text{AI}}(t) = P^{\text{main}}(t) + \sum_{j=1}^{N_{\text{tr}}} P_j^{\text{tr}}(t) + \sum_{k=1}^{N_{\text{ft}}} P_k^{\text{ft}}(t). \quad (7)$$

where P^{main} is the power of the dominant training run, P_j^{tr} that of the j -th of N_{tr} smaller training runs, and P_k^{ft} that of the k -th of N_{ft} fine-tuning runs. Each term is an independent draw of Equation (2), with its own cadence, phase ratios, and level offsets. P^{main} and the P_j^{tr} use the training parameters and the P_k^{ft} use the fine-tuning parameters; only P^{main} carries the cadence drift of Equation (3). Peak power is split between the dominant run, the small trainings, and the fine-tunings in the ratio $w_{\text{main}} : w_{\text{tr}} : w_{\text{ft}}$, and each background share is divided equally among its members.

	Symbol	Train	Fine-tune
Cadence band [Hz]	$[f^{\text{lo}}, f^{\text{hi}}]$	[0.5, 1.5]	[0.3, 0.7]
Period jitter	σ_ξ	0.1	0.1
Phase ratio	$[\rho^{\text{lo}}, \rho^{\text{hi}}]$	[0.55, 0.8]	[0.7, 0.9]
Offset spread	σ_Δ	0.05	0.03
Down-phase offset	μ_Δ	0.3	0.8
Up noise range	s^{up}	[0.02, 0.05]	[0.01, 0.03]
Down noise range	s^{dn}	[0.01, 0.03]	[0.005, 0.02]
Noise scale-down	κ_ν	0.15	0.15
<i>Cadence drift (dominant run only)</i>			
Reversion [iter ⁻¹]	θ	0.02	
Excursion std [Hz]	σ_f	0 (swept)	
Drift floor	g_{min}	0.1	
<i>Aggregate</i>			
Background counts	$N_{\text{tr}}, N_{\text{ft}}$	4, 4	
Power shares	$w_{\text{main}}:w_{\text{tr}}:w_{\text{ft}}$	9 : 0.5 : 0.5	
<i>Inference null</i>			
Mean serving load	B	0.5	
Decode RMS	A	0.40	
Decode band [Hz]	$[f_b^{\text{lo}}, f_b^{\text{hi}}]$	[0.1, 4]	
MoE corr. time [s]	τ_u	5	
MoE depth	m	0.4	
Prefill rate [Hz]	λ_p	0.08	
Prefill height	a	0.15	
Prefill duration [s]	d_p	0.1	

Table 1. Workload-model parameters of Section 2. Training and fine-tuning values are Ko and Zhu’s (Ko & Zhu, 2025), except κ_ν , the drift parameters, and σ_ξ for fine-tuning (Ko and Zhu list it for training only; we reuse it). The inference null is ours. Ko and Zhu’s intra-phase noise is a sub-millisecond fluctuation, and each sample of our tens-of-hertz generator averages over many of them, so we scale its standard deviation by κ_ν ; the s ranges are Ko and Zhu’s values before scaling. Amplitudes are fractions of the peak power P_{peak} . The drift excursion σ_f is swept in Section 5.

Each small training draws its cadence from the same band as the dominant run, so the background carries weak lines of its own inside the search band.

We define the null aggregate as the same datacentre with the dominant training run replaced by inference (Section 2.3). The background is unchanged,

$$P_{\text{AI}}^{\text{null}}(t) = P^{\text{inf}}(t) + \sum_{j=1}^{N_{\text{tr}}} P_j^{\text{tr}}(t) + \sum_{k=1}^{N_{\text{ft}}} P_k^{\text{ft}}(t), \quad (8)$$

with P^{inf} from Equation (6) at the same share w_{main} . The two aggregate classes are then power-matched and both contain the background’s weak lines. Figure 1 shows the four resulting workload classes as a verifier would see them.

3. Threat and measurement model

This section fixes what the verifier sees and who is trying to hide from it. The question a governance regime needs answered is whether a monitored facility is running training

it has not declared. For a monitored cluster the declared workload is overwhelmingly inference, so the population against which training must be detected is the inference null of Section 2.3. We consider two provers: a **benign** (non-adaptive) prover who runs its workload without regard to the meter, and an **adaptive** (hiding) prover who knows the detector and reshapes its schedule to evade it. Most of what is interesting lies in the gap between them.

3.1. The measurement model

The verifier holds one scalar time series from a meter outside the prover’s control. Writing $P_{\text{fac}}(t)$ for the total power drawn behind that meter — the static overhead P_0 plus either a single workload of Section 2 or the aggregate of Equation (7) — the meter reads

$$P_{\text{meter}}(t) = (h * P_{\text{fac}})(t) + \eta(t), \quad (9)$$

where h collects power-delivery dynamics, meter integration, reporting filters, and sampling, and $\eta(t)$ is measurement noise the prover cannot forge. Co-resident load needs no separate term: it is the background of the aggregate, and it passes through the same channel as the dominant workload. The linear time-invariant form of h is an explicit first approximation; h may be operating-point dependent. Until simultaneous on-device and external-meter measurement is performed, claims about what reaches an off-chip sensor are predictions (Section 7.2).

The implemented axes of h are a Butterworth low-pass standing for power-delivery and reporting bandwidth; a trailing boxcar integration window; zero-order-hold decimation onto a slower sample grid, deliberately without an anti-alias filter so that aliasing is representable; an in-band notch of tunable depth and Q for a transfer-function anti-resonance near the cadence; and a stable gain. The additive terms are white or AR(1) meter noise, Ornstein–Uhlenbeck baseline wander, and a square-wave controller limit cycle whose instantaneous frequency itself wanders. One honest limitation: the realised quantisation is temporal — the integration window and the sample grid — as the model has no amplitude-quantisation knob.

The default configuration is an exact identity: with every axis off, the map returns P_{fac} unchanged and consumes no randomness. That is what makes a one-axis-at-a-time sweep interpretable, since any movement is attributable to the axis being swept. Figure 2 runs that sweep on a fixed honest training workload, reporting how much of the cadence survives as the ratio of spectral power at the known f_0 to the in-band power the meter itself contributes. All six axes degrade it, and two extinguish it outright within the range a real instrument might occupy: a 1 s integration window and a 1 Hz sample rate. This is a statement about the signal in the channel, not about any detector; what each detec-




Figure 1. The four workload classes at the power meter, on the nominal channel (20 Hz sampling, no filtering or decimation, $\sigma_\eta = 4$ W). Top: a 20 s excerpt. Bottom: the multitaper spectrum over the 0.3–1.7 Hz search band, with the true cadence marked. Training, fine-tuning, and the aggregate carry a line at f_0 ; the inference null does not. Cadences are pinned for legibility ($f_0 = 0.9$ Hz training, 0.5 Hz fine-tuning); elsewhere they are drawn from the Ko bands.

tor class does with what survives, and hence the minimum meter specification, is Section 5.2.

3.2. The adversary: operators and budgets

The verifier sees power, but the prover’s costs are counted in computation. The two are linked by

$$P_{\text{fac}}(t) = r(t)F(t) + P_0(t), \quad (10)$$

where $F(t)$ is the operation rate, $r(t)$ the energy cost per operation, and $P_0(t)$ the non-compute overhead. In all our simulations r and P_0 are constant, so the power waveforms of Section 2 are operation-rate waveforms up to an offset and a scale. The distinction still matters, for two reasons. First, the cost of hiding is paid in F : an attack that reshapes the power trace by idling or reorganising work lowers throughput, and Section 6 measures that loss. Second, the meter constrains only the product $r(t)F(t)$, so a different operation rate with a compensating energy cost per operation produces the same trace; Section 7.1 returns to this.

The adaptive prover’s objective is *concealment*: to train while the trace reaching the meter stays consistent with the null. To that end it may apply, singly or in composition:

- **phase attack**: jitter, drift, or randomised work per iteration;
- **frequency relocation**: move the cadence toward or outside the search family;
- **additive burial**: independent workload superposition / spectral-floor jamming;

- **amplitude shaping**: added load flattening the compute/communication swing;
- **dilution**: reduce the training duty fraction among other work.

With unlimited superposition, dilution, shaping power, time, or freedom to restructure work, the prover can drive the signal-to-noise ratio arbitrarily low or synthesise the null; robustness claims are therefore made only over specified attack families at specified budgets.

We implement these operators inside the generator of Section 2, so that each attack costs what the physics says it costs rather than being applied as a post-hoc filter. *Work variation* quantises the up-phase into an integer number $G_i \geq 1$ of gradient-accumulation micro-steps of duration $t_\mu = \rho T_{\text{base}}/G_0$, so that $T_i^{\text{up}} = G_i t_\mu$, and varies G_i about its honest value $G_0 = 8$; crucially the down-phase keeps its honest draw, because a synchronous exchange does not shrink because the preceding compute phase was shorter. This moves real work rather than inserting idle time, and it is the reason the attack is nearly free. *Phase slip* diffuses iteration boundaries by a Brownian increment while holding the target period fixed — distinct again from ξ_i , which ties boundary diffusion to the period variance. *Harmonic smoothing* convolves the level waveform of Equation (2) with a raised cosine, ramping the transitions so the harmonic comb dies while the fundamental survives. *Shape filling* raises the down-phase level toward the up-phase, collapsing the swing at matched mean power. Section 6 prices each of these against measured hardware anchors.




Figure 2. Sensitivity of the cadence signature to the observation map, one axis at a time on a fixed honest training workload ($f_0 = 0.9$ Hz, 20 realisations, median and interquartile range). The measure is spectral power at the known f_0 over the in-band power the meter itself contributes, normalised to the nominal channel, so that attenuating axes act on the numerator and additive ones on the denominator and every panel reads the same way: down is worse for the verifier. Crosses mark settings where the cadence is not observable at all. Note f_0 is pinned away from 1.0 Hz, which would coincide with both the 2 Hz Nyquist frequency and the first sinc null of a 1 s boxcar. Reproduced by `scripts/plot_scenario_meter.py`.

4. Detection method

The question is whether the trace carries the iteration cadence of synchronous training — a narrow spectral line at f_0 , with its harmonics — that the inference null does not. A benign prover’s cadence is stationary, and a stationary line has classical tests. But a real run’s cadence drifts even without an adversary (Section 2.1), and an adaptive prover can make it wander further or vary the work per iteration, which breaks a fixed-frequency test cheaply (Section 6). The detector therefore has to follow the line. A detection establishes coherent repeated power structure, not training semantics; what it does and does not certify is taken up in Section 7.1.

4.1. Viterbi line tracking

We follow the line with the hidden-Markov frequency tracking developed for passive sonar (Streit & Barrett, 1990) and carried into continuous-gravitational-wave searches (Suvorova et al., 2016; Bayley et al., 2019; Djurović & Stanković, 2011). The trace is cut into a spectrogram over the search band 0.3–1.7 Hz, with 16 s windows overlapping by half, and each time frame j is normalised by its median over the band, giving a log-contrast L_{kj} in frequency bin k . The hidden state is the frequency bin; the prior that a real cadence wanders slowly enters as a penalty λ per bin of frequency change between frames. The detection statistic is

the best path score, averaged over the N_t frames,

$$V = \frac{1}{N_t} \max_{k_1, \dots, k_{N_t}} \sum_{j=1}^{N_t} [L_{k_j j} - \lambda |k_j - k_{j-1}|], \quad (11)$$

with no penalty on the first frame and $\lambda = 1$ throughout. Dynamic programming evaluates the maximum over all paths in time linear in N_t . A drifting line keeps a high score because the path bends to follow it. Broadband power — the continuous-batching floor of the inference null — does not, because no single slowly varying path stays on top of it, and the per-frame median normalisation keeps the null’s raised broadband level from masquerading as a line. The recovered path is also an estimate $\hat{f}_0(t)$ of the instantaneous cadence, which Appendix B uses to de-warp the trace.

Setting the threshold. V has no analytic null distribution: the maximisation over paths is a search, with the look-elsewhere effect that implies. We set thresholds empirically, at the quantile of V over the null population that gives the target false-alarm rate. This calibrates V against the modelled inference null. A deployed verifier has to obtain the same threshold from data it holds: from background traces known to be free of training, as continuous-wave searches set their tracking thresholds from noise-only realisations (Suvorova et al., 2016), or from a per-trace surrogate null of the kind we construct for the order-tracked statistic in

Appendix B. We have not yet run V through that surrogate calibration.

Instrument lines. V answers whether a persistent, slowly wandering line is present, not what produced it. A periodic physical confound — the power-controller oscillation measured on an A100 (Section 7.2) — is such a line, and against a population containing it the tracker has no power (Appendix C). Continuous-wave searches meet the same problem in instrumental lines, and handle it with a veto: lines catalogued from detector characterisation data are excluded, and candidates coincident with them discarded. The analogue here is to extract several candidate paths from the spectrogram and discard those consistent with the facility’s catalogued lines. We state it as the deployment procedure but have not evaluated it; its limits are discussed in Section 7.1. Appendix B reports a waveform-level alternative, an order-tracked cyclostationary test, that partly separates a modelled controller from training.

Comparators. Two fixed-frequency tests serve as baselines: a spectral matched filter (peak power spectral density at the searched cadence) and the Thomson multitaper harmonic F-test (Thomson, 1982; Percival & Walden, 1993). Both are honest tests for a stationary line. The multitaper test’s analytic search correction proves anti-conservative by a factor of 3.6–6, so it too is thresholded empirically; Appendix A gives the details.

5. Detection results

We report two things: the tracker’s power against the inference null as the cadence wanders, measured against an optimality ceiling, and the minimum meter that result requires.

TODO: every result in this section compares a *single* training run against the *single* inference null, with no background load (the NP ceiling and drift figure call `training_F` directly; the meter boundary uses the single-workload generator). The multi-workload aggregate of Equation (7) appears only in the dilution attack and the co-resident control. We may want to make the hard case the headline: a dominant training run over small trainings and fine-tunings, against the inference-dominant aggregate null, so the tracker must find the *dominant* line among weak background lines.

5.1. Power under cadence wander

We fixed the detector set, the conditions, and the metric before scoring any signal, and compared the tracker with the two fixed-frequency comparators on a cadence-drift axis $f_0^{\text{drift}} \in \{0, 0.1, 0.2, 0.4, 0.8, 1.5\}$ Hz, against the modelled inference null, with $n = 200$ traces of 300 s per class and thresholds set at an empirical false-alarm rate of 10^{-2} .




Figure 3. Detection rate at empirical false-alarm rate 10^{-2} against the inference null, versus cadence drift ($n = 200$ per class, 300 s traces). The Viterbi tracker holds the Whittle Neyman-Pearson ceiling across the wander axis; the spectral matched filter collapses once the line moves and the multitaper F-test has no usable power. The full bake-off, including the order-tracked statistics and structural-confuser negatives, is in Appendix C. Reproduced by `scripts/plot_drift_ceiling.py`.

Because we own the generators we can also benchmark against an optimality ceiling. No closed-form likelihood exists — both classes are black-box samplers — so we use a Whittle spectral likelihood ratio on the in-band periodogram, with a Monte-Carlo training template bank marginalised over cadence and fitted on a disjoint stream (Appendix C). The ceiling is perfect: detection rate 1.0 at every drift. The training line and the line-free null are separable in principle from a 300 s trace, so any shortfall is the detector’s.

Figure 3 shows the result. The Viterbi tracker sits on the ceiling across the whole wander axis, with detection rate 1.00 up to 0.8 Hz of drift and 0.98 at 1.5 Hz: against the inference null it is essentially Neyman-Pearson optimal. The fixed tests are not. The spectral matched filter detects a stationary line (0.94 at zero drift) but falls to 0.53 at 0.1 Hz and 0.21 at 0.4 Hz, and the multitaper F-test has essentially no power at empirical thresholds at any drift (≤ 0.02). Since a real run’s cadence drifts without any adversary, tracking is needed before an attack is even mounted. The one caveat on the ceiling runs in our favour: it discards harmonic-phase coherence and non-Gaussian structure, so the true optimum can only exceed it — which leaves the tracker’s position unchanged and the comparators further below.

Against periodic physical confounds rather than the inference null the picture is different. The tracker scores a power-controller limit cycle as highly as training and has no power against a population containing one (Appendix C); Section 7.1 takes up what that means.




Figure 4. Minimum meter specification: Viterbi detection rate at nominal false-alarm rate 0.05 over sample rate $\times$ integration window, honest training against the inference null ($n = 200$ per class). The live region is sample rate ≥ 2 Hz and integration window ≤ 0.5 s. The per-detector grid is in Appendix E. Reproduced by `scripts/plot_st2_meter_boundary.py`.

5.2. A minimum meter specification

The detection result assumes a meter that resolves the cadence band. The verifier’s instrument is itself an attack surface: a prover who cannot de-periodicise cheaply may instead argue for, or be handed, a meter too coarse to see the line. We therefore sweep the meter rather than the schedule — sample rate and integration window between the nominal 20 Hz channel and a 1 Hz integrating sampler, plus transfer-function notch depth and position near the cadence — with honest training against the hard inference null, and locate where each detector class dies.¹ The result is a minimum specification the channel must meet for detection to hold, which we state as a channel requirement, not an attack cost.

The specification. The Viterbi tracker holds full power (detection rate $\gtrsim 0.9$ at nominal 0.05) if and only if the meter satisfies three conditions together (Figure 4). First, a sample rate of at least about 2 Hz: rates of 20, 10, 5, and 2 Hz all hold the Viterbi tracker near 1.0, 1 Hz collapses it to 0.36 as the Nyquist edge crosses the band, and 0.5 Hz is dead (every detector at chance). Second, an integration window no longer than about 0.5 s: at 20 Hz sampling a window of 0.5 s still holds (0.99), but a 1 s boxcar sinc-nulls the 1 Hz band centre and collapses the tracker to 0.25 even at full sample rate. Third, no deep transfer-function notch at the cadence: an off-band or shallow notch is tolerable (tracker ≈ 1.0), but a full-depth notch at 1.0–1.2 Hz roughly halves tracker power (to 0.48–0.55) without erasing it (Figure 14).

¹Frozen on the MATS compute cluster (slurm array 5498), 48 cells, $n = 200$, 300 s traces; `results/st2/meter_boundary_summary.json`. Marginal and dead-cell AUCs are sensitive to the BLAS build, but the live-region boundary is not.

6. The de-periodicisation frontier

The detection results assume the prover leaves a recoverable cadence. An adaptive prover need not: it can de-periodicise its schedule to smear that cadence away. The question this section answers is not whether de-periodicisation is possible — it always is — but how far a prover must go before the meter loses the signal, and what that costs. This is a central result of the paper, not a caveat, and its shape is not the naive one. De-periodicising is not automatically cheap: some ways of doing it cost throughput without buying any concealment, and the interesting question is whether *any* way is both cheap and effective.

We sweep six attack families as operators on the training schedule (Section 3.2) — per-iteration duration jitter, slowly drifting cadence, variable real work per iteration, cadence relocation toward the search-band edge, harmonic suppression, and dilution by co-resident load — plus a meter-degradation family treated separately in Section 5.2. For each family we step an attack level and, at every level, measure the physical de-periodicisation it induces (cadence coefficient of variation, phase diffusion) against the power of the Viterbi tracker and of the fixed tests (the spectral matched filter and the multitaper F-test, taking the better of the two at each cell), all at nominal false-alarm rate 0.05 with $n = 200$ traces per class. The order-tracked statistics of Appendix B were scored on the same cells and are reported with the full sweep in Appendix D; on every cell with a measured cost they add nothing to the tracker. Where an attack has a measured systems cost we attach it from the single-A100 throughput anchors (Section 7.2); the plausible *learning-efficiency* penalty — batch-size, accumulation, or update-timing changes that preserve tokens per second while worsening progress per token or per joule — is not something a synthetic power model can measure, and we carry it as the standing open empirical question rather than a result.²

The costly attacks are the wrong ones. The de-periodicisations that degrade the line most are exactly the ones a rational prover would not pay for. Idle-insertion jitter buys no concealment: it costs 15–375% throughput and leaves the tracked line detectable. Slowly drifting the cadence costs up to $\sim 140\%$ and is precisely the regime the tracker was built to survive (Section 5.1). Power shaping costs 170–230%. Against every one of these the tracker holds while the prover pays a throughput tax that a compute-governance regime would itself flag. Measured work on real

²Frozen on the MATS compute cluster (slurm array 5635), $n = 200$ per class over 50 cells, seed 0, the full five-detector set; `results/st2/frontier_summary.json` (`provisional=false`). The systems-cost axis reads only the clean throughput-overhead field of the measured campaign, not its signature-B line-band arrays, so the frontier is unaffected by the transport caveat of Section 7.2.

hardware points the same way from a different channel: four evasion strategies run on a single H200 — chunked optimizer updates, dilution with decoy inference, throttling, and diluted LoRA — cost between 28 and 69% of throughput (Gargiulo & Kulp, 2026). Each of them inserts or substitutes work, which is what makes it expensive, and none varies the real computation per iteration — the family we turn to next.

The cheap attack collapses the fixed tests but not the tracker. The dangerous family is variable real work. Varying the useful computation performed per iteration moves the cadence without inserting idle time, so its throughput cost is essentially zero: at work-variation level 0.35 and 0.5 the measured overhead is -0.10% and -0.36% — zero within noise. At that price the fixed tests are gone. The spectral matched filter falls to detection rate 0.26 and 0.27 (its AUC collapses to 0.44 and 0.41, below chance), and the multitaper F-test sits at its floor. Yet the Viterbi tracker holds detection rate 1.0 at both levels (Figure 5; per-detector curves in Appendix D). This is the genuine scope boundary of the passive channel: a prover willing to vary real work evades every cheap, fixed detector at no measurable systems cost, and only a tracking detector still sees the schedule.

The tracker’s own boundary is unmeasured. The Viterbi tracker does eventually bend — at work-variation level 0.7 its detection rate falls to 0.59 — but we do not report that as a defeat, because we cannot price it. The throughput cost of driving work variation that hard is not in the measured anchors, and the learning-efficiency cost is the open question above. The honest statement is therefore asymmetric: within the regime whose cost we can measure (overhead within noise) the tracker wins outright, and the far end of the frontier, where the tracker finally fails, is a cost we could not quantify rather than a cost shown to be small. That the tracker’s advantage here is real and not a tuning artefact is established independently in Section 5.1: against the inference null the Viterbi tracker is essentially Neyman–Pearson optimal, so the fixed tests’ collapse under variable work is a genuine loss of information from the schedule, not a deficiency of a particular detector.

The trade-off itself. Plotting hiding directly against measured cost (Figure 5) separates the two detectors sharply. Against the fixed tests the Pareto staircase reaches hiding 0.74 at essentially zero overhead and climbs only to 0.79 at 159% and 0.82 at 375% however much the prover spends: the cheapest attack is within 0.08 of the best one. Against the Viterbi tracker the staircase never leaves the floor, rising only to 0.16 and doing so at 159% overhead, so no attack in the anchored set buys meaningful concealment from the tracker at any price we can measure. The two staircases together are the paper’s central empirical claim in one figure: what the fixed tests lose for free, the tracker keeps at




Figure 5. Cost of hiding. Measured throughput overhead against hiding (1 — detection rate at false-alarm rate 0.05) for the four attack families with a hardware cost anchor, with the Pareto staircase — the best hiding available at or below a given cost — against the fixed tests (orange; the better of the matched filter and multitaper F-test per cell) and the Viterbi tracker (blue). Varying real work per iteration hides ~ 0.74 from the fixed tests at no measurable cost; against the tracker no anchored attack hides more than 0.16. Cells without a cost anchor are shown as ticks in the hatched strip, among them work-variation level 0.7, where the tracker begins to fail. Reproduced by `scripts/plot_money_pareto.py`.

every measured cost. The figure also shows the boundary of that claim honestly. Only 16 of the 50 swept cells have a measured hardware cost anchor, and the rest — including work-variation level 0.7, where the tracker finally bends — carry no price at all, so they appear in a separate unpriced strip rather than being dropped.

A lower bound on the price of hiding. The frontier bounds hiding from above: it says what a prover can hide from our detectors, which a better detector would move. The complementary direction — what hiding must cost — can be argued for one attack family (Section G.1). For i.i.d. period jitter of scale σ , no test’s detection advantage exceeds the total variation between the training and null laws, so a verifier that demonstrably attains advantage $J(\sigma) > \varepsilon$ rules out ε -covertiness at that σ . With the simplest such verifier, a fixed-bin threshold test at the nominal cadence, ε -covertiness requires $\sigma^*(0.5) = 0.18$ and $\sigma^*(0.2) = 0.36$, which the measured anchors price at tens of percent of throughput and at order 100% respectively. Three limits apply. The bound is read off the weakest verifier, so it is loose against the tracker, which Figure 5 shows to be stronger; it covers one attack family; and turning distortion into cost rests on an assumption — that at fixed useful work, jitter costs throughput or learning efficiency — rather than a result.

7. Discussion and conclusion

We set out to build a method that turns a low-rate external power trace into a test for iteration-structured training, to price its evasion against a prover who knows the test, and to state what the channel cannot certify at any price. The method is a Viterbi line tracker. Against the inference null that stands for a facility’s declared workload it is essentially Neyman–Pearson optimal across cadence wander (Section 5.1), provided the meter samples at ≥ 2 Hz with ≤ 0.5 s integration (Section 5.2). Varying real work per iteration evades every fixed-frequency test at no measurable throughput cost, but no cost-anchored attack buys the prover more than 0.16 hiding from the tracker (Section 6). This section says what such a detection certifies, where the synthetic evidence stops, and how the result reads for a regulator.

7.1. What a detection certifies

A detection establishes that the trace carries a training-shaped periodic schedule. It does not establish that the computation is training. We built seven control loads whose training status we know and scored them with three decision rules fitted to separate training from the inference null (Appendix F). Five non-training loads — a training-shaped non-ML kernel loop, a controller limit cycle, a periodic inference request generator, a gradient-only computation, and a discarded-update decoy that runs the full forward, backward and optimiser step and throws the update away — all score as training under every rule (Table 2).

The adversary operators of Section 3.2 deliberately omit the extremal member of this class. A *semantic decoy* — a load reproducing the physical schedule of training without retaining a trained model — is not an attack a concealing prover would mount, since manufacturing evidence of the activity it is hiding buys it nothing. It enters this paper instead as a non-training control (Table 2) and as the construction behind the non-identifiability argument below, and it carries weight for two reasons that need no motivated decoy-runner: its benign counterparts occur in ordinary datacentre operation, and the active protocol of Appendix H inverts the prover’s objective, at which point a *challenge-aware decoy* — a correlated dummy response mimicking the prescribed reaction without training — does become a live attack.

The ceiling is a property of the channel, not of our classifiers. Write w for a computational execution, h for the hardware and power-delivery path, and $K_h(w)$ for the law of the observed trace that Equations (9) and (10) induce. Call two executions *observationally equivalent* when $K_{h_0}(w_0) = K_{h_1}(w_1)$. A passive verifier sees only the trace, so its decision is some measurable function of it, and on an observationally equivalent pair every such function fires with the same probability under both. No passive scalar-power detector can therefore certify the semantic content

of a computation against an unrestricted adversary. A non-vacuous guarantee has to put the equivalent pairs out of scope, and there are four ways to do that within the passive channel — restricting the admissible non-training workload class, the adversary’s freedom to add or restructure computation, the hardware and power-transfer family, or the physical property claimed — alongside a fifth route that leaves the passive channel altogether by changing the protocol (Appendix H). We do not claim the list is exhaustive.

The construction behind this is nearly available on the prover’s own hardware. A program can execute the same forward, backward, and optimiser kernels as training and discard each updated state, making its trace arbitrarily close to genuine training while its semantic classification differs. This is not a hypothetical weakness; it is what Table 2 measures. That outcome is a property of the channel rather than of our classifiers, which is why this ceiling cannot be removed by a better detector and why escaping it requires an active challenge (Appendix H).

The argument binds even where no adversary is present. A concealing prover has no reason to run a decoy (Section 3.2), but most of the equivalence class is not adversarial at all: a training-shaped non-ML kernel loop, a controller limit cycle, and a periodic inference request generator all sit inside it, and all arise in ordinary operation. The consequence is two-sided. A rule phrased semantically — training rather than serving, which model, whose weights — cannot be adjudicated from the schedule alone in either direction, so the same equivalence that would let a decoy counterfeit training also caps the evidentiary weight a verifier can attach to a positive detection.

Instrument lines and false alarms. The same reading covers the tracker’s one systematic failure. A periodic physical confound, such as the power-controller oscillation measured on an A100 (Section 7.2), is a persistent wandering line, and the Viterbi score cannot tell it from training (Appendix C). The failure matters for honest operators rather than concealing ones: a controller line does not help a prover hide training, but it can flag a compliant inference facility. Three things bound it. First, a known-line veto (Section 4.1) discards candidate paths that coincide with lines catalogued from the facility’s own characterisation data. We have not evaluated it, and it carries a cost a gravitational-wave veto does not: the adversary knows the veto list, and can relocate its cadence into a vetoed band. Second, per-card controllers need not add coherently at a facility meter (Section 7.2). Third, a positive detection is a trigger for escalation rather than a verdict, and the active protocol of Appendix H resolves a controller false alarm as it resolves the decoys above. The order-tracked cyclostationary statistic of Appendix B partly separates the modelled controller from training by waveform shape, but only conditionally, and we do

not rely on it.

7.2. Reality check and transfer limits

Every result above is synthetic. The one place we hold the scenario model against measured hardware, it fails in an instructive way, and we report that failure as a result rather than a closing nod.

The negative transport case. On measured single-A100 traces the iteration cadence is present but shallow: a saturated card has no communication phase, so the iteration structure modulates a ~ 370 W draw by only a few watts ($\sim 0.4\%$ of board power). Worse for the naive reading, the spectra contain a 0.31–0.45 Hz component, roughly ten times stronger than the training line, that appears whenever the load varies in time but not at idle or under constant load — the signature of the card’s own power-management loop. It is a hardware line that says “time-varying load,” not “training,” and it is the instrument line of Section 7.1: the standing counterexample to reading any detected line as training.

Why this is the honest outcome. Two things make the single-card signature shallow. One saturated card has no synchronous communication barrier, so the deep compute–communication contrast the scenario model assumes is absent; and the generator is literature-parameterised, not calibrated to the measured channel, a gap of about $75\times$ in modulation depth. The synthetic results are therefore a sensitivity analysis over a plausible parameter family, not a prediction calibrated to this hardware — which is why the paper’s claims are scoped to the stated generator and its stated nulls throughout.

The single-card gap is shared, and lands differently. The measured side-channel corpus that separates training from inference at high bandwidth is also a single card, with cross-device generalisation left open (Gargiulo & Kulp, 2026). What differs is which observable the gap touches. A board-level probe reads the optimizer-update phase, which is present on one card, so the absent distributed setting costs that result little. The cadence is a collective effect of synchronisation, so a single card is exactly the case in which it is weakest — which is why the same gap is a transport caveat there and a first-order limitation here.

Falsifiable predictions. The model makes three predictions we cannot test here but state as falsifiable. Distributed synchronisation should deepen the line: at cluster scale the periodic all-reduce barrier produces a deep, coherent, legible cadence rather than the shallow single-card ripple, which is precisely the physics the cost-of-hiding assumption leans on. And an external wall or rack meter should see a different,

lower-passed trace than on-board NVML telemetry, shifting where the minimum meter specification of Section 5.2 binds. And per-card power-controller oscillations, which dominate the single-card spectrum, should add incoherently across a cluster — roughly as $\sqrt{N}$ for N independent controllers, against N for a synchronous training line — and so fall relative to the training line at a facility meter. All three are predictions, not results.

7.3. The claim ladder

The results can be read back as a *claim ladder* (Table 3): four claims of increasing governance relevance, each bought at a stated cost in verifier information, and each with a hard ceiling that motivates the next. Read as a procedure, it is screening followed by escalation. Rung 1 establishes that coherent periodic power structure is present, which is evidence rather than identification. Rung 2 converts that structure into training-likeness at the price of committing to an inference null; the detection results of this paper live here, and its ceiling is semantic non-identifiability. Rungs 3 and 4 leave the passive channel for an active protocol and a transcript primitive, and are specified but not run (Appendix H). The ladder is a ladder of claims and assumptions, not a quantitative bound on unidentified compute.

7.4. Scope, limitations, and governance reading

Scope. Three restrictions fix what the results cover. First, the observable: we use the time-resolved power trace and exclude the integrated-energy channel, so there is no bounding of how much compute was performed and no energy-per-token pricing as a verification instrument. This is a restriction on the observable, not a claim that everything lives in the spectrum; cyclostationary structure is expressed spectrally, but the challenge correlation of the active protocol is a time-domain operation on the same trace. Second, the population: we do not claim a recoverable cadence is universally necessary for training, and some conceivable execution can be made fully aperiodic. We target efficient, synchronous, iteration-structured training that exposes a recoverable cadence, the expected physical signature of that regime where distributed synchronisation makes the modulation deep and coherent. Destroying that cadence is then an adversarial action with a measurable budget, not a counterexample to the method, and measuring that budget is what Section 6 does. The declared workload need not be pure inference either: a facility may be permitted fine-tuning or training below a threshold, whose lines our detector would also find, and telling which line and how large a run it represents needs the integrated-energy channel we exclude. Third, the horizon: the results are fixed-horizon, and we make no claim that our statistics support anytime-valid monitoring, deferring sequential e-process treatment. The cryptographic binding of the active protocol is assumed as an ideal functionality, not

Table 2. Semantic-control outcomes: fraction of each load scored as training at a false-positive rate of 0.05, per decision rule. Five non-training loads all score as training; the only control the tracked rules degrade on is asynchronous (de-periodicised) genuine training, which they nonetheless mostly catch.

Control load	Training?	Physics	Fitted	Learned
Discarded-update decoy	no	1.00	1.00	1.00
Gradient-only	no	1.00	1.00	1.00
Non-ML kernel loop	no	1.00	1.00	1.00
Controller cycle	no	1.00	1.00	1.00
Periodic inference	no	1.00	1.00	1.00
Co-resident mixture	yes	1.00	1.00	1.00
Asynchronous training	yes	0.72	0.55	1.00

built. Real distributed hardware enters only as the falsifiable predictions and the negative transport case of Section 7.2.

Limitations. The scenario model is literature-parameterised, not calibrated to measured hardware, so every quantitative result is a sensitivity analysis over a stated parameter family; the $\sim 75\times$ modulation gap of Section 7.2 is the standing reminder. The observation map is a linear time-invariant first approximation. The optimality ceiling of Appendix C is optimal only under the Whittle model, hence a lower bound on the true optimum. Every robustness claim is bounded by the attack families and budgets we specified; a family we did not model is outside it. The cost bound of Section G.1 shares that limit: it covers i.i.d. period jitter, and while it holds against any verifier class containing the test it exhibits, it is obtained from a fixed-bin test and so reads loose against the Viterbi tracker our own experiments show to be stronger — it does not yet close the frontier’s far end. Its cost leg is an assumption, not a result, for the reasons set out in Section G.1. The tracker’s thresholds are calibrated by Monte Carlo on the modelled null, not per trace, and the known-line veto that would handle instrument lines is specified but not evaluated. And the learning-efficiency leg of that bound — the axis that would turn it into a single priced statement rather than a disjunction — is argued, not measured, and remains the central open empirical question.

Governance reading. What a regulator can buy today is a screen. From an external meter alone, meeting the specification of Section 5.2, a Viterbi tracker detects the iteration cadence of synchronous training against a declared inference workload, robust to cadence wander, drift, and the zero-cost variable-work evasion that defeats every fixed-frequency test; within the attacks whose cost we could measure, none buys meaningful concealment from it. It does not certify that the computation is training rather than a schedule-matched mimic, benign or deliberate, and it cannot by itself tell training from a periodic instrument line. Both are limits of the passive channel, not of the method, and they run both ways: they cap what a positive detection proves as well as what a counterfeit could fake. A positive detection should therefore

trigger escalation — to the active protocol of Appendix H and the cryptographic primitive it assumes — rather than stand as proof of training. The passive meter is a cheap, deployable front line that narrows the question and prices the cheap evasions.

References

- Baker, M., Kulp, G., Marks, O., Brundage, M., and Heim, L. Verifying international agreements on AI: Six layers of verification for rules on large-scale AI development and deployment, July 2025. URL <https://arxiv.org/abs/2507.15916>.
- Bash, B. A., Goeckel, D., and Towsley, D. Limits of reliable communication with low probability of detection on AWGN channels. *IEEE Journal on Selected Areas in Communications*, 31(9):1921–1930, 2013. doi: 10.1109/JSAC.2013.130923.
- Bayley, J., Messenger, C., and Woan, G. Generalized application of the Viterbi algorithm to searches for continuous gravitational-wave signals. *Physical Review D*, 100:023006, 2019. doi: 10.1103/PhysRevD.100.023006.
- Bonnardot, F., El Badaoui, M., Randall, R. B., Danière, J., and Guillet, F. Use of the acceleration signal of a gearbox in order to perform angular resampling (with limited speed fluctuation). *Mechanical Systems and Signal Processing*, 19(4):766–785, 2005. doi: 10.1016/j.ymssp.2004.05.001.
- Brands, S. and Chaum, D. Distance-bounding protocols. In *Advances in Cryptology — EUROCRYPT ’93*, volume 765 of *Lecture Notes in Computer Science*, pp. 344–359. Springer, 1994. doi: 10.1007/3-540-48285-7_30.
- Cachin, C. An information-theoretic model for steganography. *Information and Computation*, 192(1):41–56, 2004. doi: 10.1016/j.ic.2004.02.003.
- Canetti, R. Universally composable security: A new paradigm for cryptographic protocols. In *Proceedings*

Table 3. The claim ladder. Each rung answers a strictly more governance-relevant question at a stated cost in verifier information, and each rung’s hard ceiling motivates the next. The fourth column attaches this paper’s evidence.

Rung	Claim	Additional requirement	re-	Evidence in this paper	Hard ceiling
1. Structural evidence	Coherent periodic power variation exists under a specified structural null	External resolved only	time-trace	Order-tracked cyclic structure is detected with surrogate-calibrated, level-safe false-alarm control on linear-stationary nulls (Appendix B); the Viterbi score alone cannot exclude a controller (Appendix C)	Periodic physical confounds; structural-null misspecification; die-to-meter attenuation
2. Conditional classification	The trace is more training-like than the stated inference population	A representative or explicitly modelled inference null		Viterbi detects training at the Neyman–Pearson ceiling across cadence wander (Section 5.1) and against zero-cost variable real work (Section 6), given a meter of ≥ 2 Hz and ≤ 0.5 s integration (Section 5.2); fitted rules reach area under the curve 1.00 under six domain shifts (Appendix F)	Domain shift; strong superposition, dilution or shaping; cadence moved outside the searched family; semantic non-identifiability (Table 2)
3. Active physical authentication	A load behind the meter responded to an unpredictable verifier challenge	Prover participation; a challenge-bound schedule		Specified, not run (Appendix H)	A challenge-aware dummy load; no semantic binding by power alone
4. Work verification	The physical response is tied to declared training state transitions	A sound, fresh transcript or proof and an explicit binding assumption		Assumed as an ideal functionality (Appendix H)	The remote-proof / colocation gap, unless the assumed primitive closes it

of the 42nd IEEE Symposium on Foundations of Computer Science (FOCS), pp. 136–145. IEEE, 2001. doi: 10.1109/SFCS.2001.959888.

Chari, S., Rao, J. R., and Rohatgi, P. Template attacks. In *Cryptographic Hardware and Embedded Systems — CHES 2002*, volume 2523 of *Lecture Notes in Computer Science*, pp. 13–28. Springer, 2002. doi: 10.1007/3-540-36400-5_3.

Clark, S. S., Ransford, B., Rahmati, A., Guineau, S., Sorber, J., Xu, W., and Fu, K. WattsUpDoc: Power side channels to nonintrusively discover untargeted malware on embedded medical devices. In *2013 USENIX Workshop on Health Information Technologies (HealthTech 13)*. USENIX Association, 2013.

Cover, T. M. and Thomas, J. A. *Elements of Information Theory*. Wiley-Interscience, Hoboken, NJ, second edition, 2006. doi: 10.1002/047174882X.

Dandawaté, A. V. and Giannakis, G. B. Statistical tests for presence of cyclostationarity. *IEEE Transactions on Signal Processing*, 42(9):2355–2369, 1994. doi: 10.1109/78.317857.

Djurović, I. and Stanković, L. Viterbi algorithm for chirp-rate and instantaneous frequency estimation. *Signal Pro-*

cessing, 91(5):1308–1314, 2011. doi: 10.1016/j.sigpro.2010.10.007.

Fisher, R. A. *The Design of Experiments*. Oliver and Boyd, Edinburgh, 1935.

Gargiulo, S. and Kulp, G. Workload identification with physical side channels for AI governance, 2026. URL <https://arxiv.org/abs/2609.00309>.

Grattafiori, A., Dubey, A., Jauhri, A., Pandey, A., Kadian, A., et al. The Llama 3 herd of models, 2024. URL <https://arxiv.org/abs/2407.21783>.

Hart, G. W. Nonintrusive appliance load monitoring. *Proceedings of the IEEE*, 80(12):1870–1891, 1992. doi: 10.1109/5.192069.

Heim, L. and Koessler, L. Training compute thresholds: Features and functions in AI regulation, 2024. URL <https://arxiv.org/abs/2405.10799>.

Jain, R., Tran, B., Chen, K., Sinclair, M. D., and Venkataraman, S. PAL: A variability-aware policy for scheduling ML workloads in GPU clusters. In *SC24: International Conference for High Performance Computing, Networking, Storage and Analysis*, pp. 1–18. IEEE, 2024. doi: 10.1109/SC41406.2024.00032.

- Jia, H., Yaghini, M., Choquette-Choo, C. A., Dullerud, N., Thudi, A., Chandrasekaran, V., and Papernot, N. Proof-of-learning: Definitions and practice. In *2021 IEEE Symposium on Security and Privacy (SP)*, pp. 1039–1056. IEEE, 2021. doi: 10.1109/SP40001.2021.00106.
- Ko, M.-S. and Zhu, H. Wide-area power system oscillations from large-scale AI workloads, August 2025. URL <https://arxiv.org/abs/2508.16457>.
- Kocher, P., Jaffe, J., and Jun, B. Differential power analysis. In *Advances in Cryptology — CRYPTO ’99*, volume 1666 of *Lecture Notes in Computer Science*, pp. 388–397. Springer, 1999. doi: 10.1007/3-540-48405-1.25.
- Kwon, W., Li, Z., Zhuang, S., Sheng, Y., Zheng, L., Yu, C. H., Gonzalez, J. E., Zhang, H., and Stoica, I. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP ’23)*, pp. 611–626. ACM, 2023. doi: 10.1145/3600006.3613165.
- Lin, J., Jiang, Z., Song, Z., Zhao, S., Yu, M., Wang, Z., Wang, C., Shi, Z., Shi, X., Jia, W., Liu, Z., Wang, S., Lin, H., Liu, X., Panda, A., and Li, J. Understanding stragglers in large model training using what-if analysis. In *19th USENIX Symposium on Operating Systems Design and Implementation (OSDI 25)*, 2025. URL <https://arxiv.org/abs/2505.05713>.
- Monfared, S. K., Ganji, F., Holcomb, D., and Tajik, S. Timing and memory telemetry on GPUs for AI governance, 2026. URL <https://arxiv.org/abs/2602.09369>.
- O’Gara, A., Kulp, G., Hodgkins, W., Petrie, J., Immler, V., Aysu, A., Basu, K., Bhasin, S., Picek, S., and Srivastava, A. Hardware-enabled mechanisms for verifying responsible AI development. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, 2025. URL <https://arxiv.org/abs/2505.03742>.
- Pappu, R., Recht, B., Taylor, J., and Gershenfeld, N. Physical one-way functions. *Science*, 297(5589):2026–2030, 2002. doi: 10.1126/science.1074376.
- Patel, P., Choukse, E., Zhang, C., Goiri, Í., Warriier, B., Mahalingam, N., and Bianchini, R. POLCA: Power oversubscription in LLM cloud providers, 2023. URL <https://arxiv.org/abs/2308.12908>.
- Patel, P., Choukse, E., Zhang, C., Goiri, Í., Warriier, B., Mahalingam, N., and Bianchini, R. Characterizing power management opportunities for LLMs in the cloud. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 3 (ASPLOS ’24)*, pp. 207–222. ACM, 2024a. doi: 10.1145/3620666.3651329.
- Patel, P., Choukse, E., Zhang, C., Shah, A., Goiri, Í., Maleki, S., and Bianchini, R. Splitwise: Efficient generative LLM inference using phase splitting. In *2024 ACM/IEEE 51st Annual International Symposium on Computer Architecture (ISCA)*, pp. 118–132, 2024b. doi: 10.1109/ISCA59077.2024.00019.
- Percival, D. B. and Walden, A. T. *Spectral Analysis for Physical Applications: Multitaper and Conventional Univariate Techniques*. Cambridge University Press, 1993. doi: 10.1017/CBO9780511622762.
- Rahman, R. and Tajdari, S. Detecting hidden ML training with zero-overhead telemetry, 2026. URL <https://arxiv.org/abs/2606.19262>.
- Sastry, G., Heim, L., Belfield, H., Anderljung, M., Brundage, M., Hazell, J., O’Keefe, C., Hadfield, G. K., Ngo, R., Pilz, K., Gor, G., Bluemke, E., Shoker, S., Egan, J., Trager, R. F., Avin, S., Weller, A., Bengio, Y., and Coyle, D. Computing power and the governance of artificial intelligence, 2024. URL <https://arxiv.org/abs/2402.08797>.
- Shavit, Y. What does it take to catch a Chinchilla? Verifying rules on large-scale neural network training via compute monitoring, 2023. URL <https://arxiv.org/abs/2303.11341>.
- Shazeer, N., Mirhoseini, A., Maziarz, K., Davis, A., Le, Q. V., Hinton, G., and Dean, J. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *5th International Conference on Learning Representations (ICLR 2017)*, 2017. URL <https://arxiv.org/abs/1701.06538>.
- Streit, R. L. and Barrett, R. F. Frequency line tracking using hidden Markov models. *IEEE Transactions on Acoustics, Speech, and Signal Processing*, 38(4):586–598, 1990. doi: 10.1109/29.52700.
- Suvorova, S., Sun, L., Melatos, A., Moran, W., and Evans, R. J. Hidden Markov model tracking of continuous gravitational waves from a neutron star with wandering spin frequency. *Physical Review D*, 93:123009, 2016. doi: 10.1103/PhysRevD.93.123009.
- The White House. Executive order 14110: Safe, secure, and trustworthy development and use of artificial intelligence. Federal Register 88 FR 75191, 2023. URL <https://www.federalregister.gov/documents/2023/11/01/2023-24283/>.
- Thomson, D. J. Spectrum estimation and harmonic analysis. *Proceedings of the IEEE*, 70(9):1055–1096, 1982. doi: 10.1109/PROC.1982.12433.

- Tsybakov, A. B. *Introduction to Nonparametric Estimation*. Springer Series in Statistics. Springer, New York, 2009. doi: 10.1007/b13794.
- Wu, T., Wang, W., Yu, Y., Yang, S., Wu, W., Duan, Q., Yang, G., Wang, J., Qu, L., and Zhang, L. FALCON: Pinpointing and mitigating stragglers for large-scale hybrid-parallel training, 2024. URL <https://arxiv.org/abs/2410.12588>.
- Yu, G.-I., Jeong, J. S., Kim, G.-W., Kim, S., and Chun, B.-G. Orca: A distributed serving system for transformer-based generative models. In *16th USENIX Symposium on Operating Systems Design and Implementation (OSDI 22)*, pp. 521–538. USENIX Association, 2022. URL <https://www.usenix.org/conference/osdi22/presentation/yu>.
- Yücek, T. and Arslan, H. A survey of spectrum sensing algorithms for cognitive radio applications. *IEEE Communications Surveys & Tutorials*, 11(1):116–130, 2009. doi: 10.1109/SURV.2009.090109.
- Zoha, A., Gluhak, A., Imran, M. A., and Rajasegarar, S. Non-intrusive load monitoring approaches for disaggregated energy sensing: A survey. *Sensors*, 12(12):16838–16866, 2012. doi: 10.3390/s121216838.

A. Fixed-frequency baselines

This appendix gives the two fixed-frequency comparators of Section 4 in full.

When the cadence is stationary, a repeated iteration structure appears as a line, or a harmonic comb, in the power spectrum, and the natural test is the Thomson multitaper harmonic F-test (Thomson, 1982; Percival & Walden, 1993). At a fixed frequency f the statistic $F(f)$ is distributed as $F(2, 2K - 2)$ under a locally smooth Gaussian spectral null with K tapers, so the test is *approximately pivotal under a stationary, locally smooth spectral null*. It is not an exact false-alarm guarantee across arbitrary external power traces: coloured baselines, controller interference, and co-resident load all violate the local-smoothness premise. Because the prover’s cadence is not known a priori, the test must search a band of candidate frequencies and harmonics, which requires a search correction.

The analytic correction is itself unreliable. Sweeping the padded frequency grid and applying a Bonferroni correction over the Rayleigh-resolution cell count is *anti-conservative by a factor of roughly 3.6–6×*, uniformly across our null suite: the maximum of F over the padded grid behaves like about five times more effective independent tests than the Rayleigh count, because the F ratio fluctuates below the Rayleigh scale. The analytic multitaper tail is therefore usable only as a ranking statistic; its operative false-alarm calibration is the Monte-Carlo one supplied by the harness of Section B.3. A spectral matched filter (peak PSD at the searched cadence) serves as the simple comparator. Both are honest baselines for a stationary line, and both fail the moment the line wanders: as the comparison of Section 5.1 shows, the fixed multitaper F-test retains essentially no usable detection power once the cadence drifts, which is exactly why the adaptive setting needs a different detector.

B. Order-tracked cyclostationary attribution and null validity

The main text detects training with the Viterbi score alone. This appendix reports a longer pipeline that the main text does not rely on: the Viterbi path defines a de-warp of the trace, the de-warped trace is tested for cyclostationarity at the tracked order, and the result is calibrated per trace against a surrogate null. Its value is attribution. Phase-coherent cyclic structure at the tracked order partly separates training from a power-controller limit cycle, which the Viterbi score cannot (Appendix C), but only when thresholds are set against a controller population, and the controller still inflates its surrogate-calibrated false-alarm rate (Section B.3).

B.1. Tracked cyclostationary detection

An adaptive prover breaks the benign baselines cheaply by letting the cycle frequency wander, smearing the line across the searched band. We restore coherence by estimating the wander and undoing it before testing. The detector has three components, each with one job. First, a tracker estimates the instantaneous cycle frequency $\hat{f}_0(t)$ from the time–frequency plane by Viterbi dynamic programming under a slow-wander transition penalty — the hidden-Markov frequency-line tracking developed for passive sonar (Streit & Barrett, 1990) and carried into continuous-gravitational-wave searches (Suvorova et al., 2016; Bayley et al., 2019; Djurović & Stanković, 2011); the tracker is a nuisance-parameter estimator, *not* the detector. Second, the phase estimate

$$\hat{\theta}(t) = 2\pi \int_0^t \hat{f}_0(s) ds \quad (12)$$

defines a resampling of the trace at equal increments of $\hat{\theta}$ (tachless order tracking (Bonnardot et al., 2005)), so a wandering line becomes a constant-order line. Third, the de-warped signal is tested for cyclostationarity at candidate cycle frequency α with the Dandawaté–Giannakis statistic (Dandawaté & Giannakis, 1994)

$$Q_\alpha = N \hat{r}_\alpha^H \hat{\Sigma}_\alpha^{-1} \hat{r}_\alpha, \quad (13)$$

where $\hat{r}_\alpha$ collects the estimated cyclic autocovariances at lags τ and $\hat{\Sigma}_\alpha$ is a consistent estimate of their long-run covariance. For a *fixed* α , Q_α is asymptotically χ^2_{2L} under the structural null, given weak dependence, consistent covariance estimation, and regularity.

Those conditions do not automatically cover the complete adaptive pipeline. The tracker performs an effective search over frequency paths (a look-elsewhere effect); the time warp is estimated from the same record it tests; the covariance estimator $\hat{\Sigma}_\alpha$ must be chosen; and the asymptotic χ^2 tail may be a poor guide at finite length. Whether the null survives these steps — and in what form it can be claimed — is the make-or-break question of Section B.3. We do not call the adaptive pipeline an analytic CFAR detector; Section B.3 establishes what it is instead.




Figure 6. Realised versus nominal false-alarm rate for the adaptive pipeline (Dandawaté–Giannakis with a Bartlett HAC covariance), one panel per stage, every null in the suite, with Clopper–Pearson 95% intervals; nulls are coloured by stage group. No stationary null exceeds the nominal rate at any stage: the two-sided miss is *deflation*, never inflation.

B.2. Surrogate calibration

Because the asymptotic tail cannot be relied on for the complete pipeline — Section B.3 shows why — we calibrate each trace against a surrogate null generated from the trace itself. Given an observed record we form S surrogates by taking its discrete Fourier transform, replacing each phase with an independent uniform draw while preserving the modulus, and inverting. The resulting series has the observed record’s power spectrum and no phase structure, so the ensemble encodes the invariance “linear and stationary with this spectrum”. The *entire* pipeline — tracker, resampling, and test — is then rerun on each surrogate, and the reported p -value is the rank of the observed statistic in that ensemble. Running the full pipeline inside the surrogate loop is the point: it is what makes the tracker’s path search and the estimated warp part of the null rather than unmodelled extra freedom.

Two properties of this construction are worth stating in advance, because they determine what the resulting test can and cannot claim. It is exact only against the invariance it encodes, so a null that is genuinely nonlinear or genuinely cyclostationary is not covered — and a controller limit cycle is exactly such a null. And it is expensive: each trace costs S additional full pipeline evaluations, which is why the surrogate campaigns below are sized smaller than the analytic ones.

B.3. Null validity of the complete pipeline

The question is whether the adaptive pipeline has a null we can defend. We answer it with a staged design that isolates each adaptive ingredient in turn, crossed with a graded suite of nulls, and measured by the realised false-alarm rate at operationally small nominal levels — not only 0.05. The four stages add one adaptive step at a time: (1) a fixed cyclic frequency α with no resampling; (2) resampling under a *known* independent time warp; (3) a warp *estimated* on held-out blocks and scored only on the complement; and (4) the fully adaptive full-trace Viterbi tracker. The null suite grades from stationary Gaussian and coloured processes (white, AR(1), a resonant AR(2)), through stationary non-Gaussian processes (heavy-tailed, skewed, squared Gaussian), to locally stationary and controller-like processes (a Van der Pol-type power-management limit cycle mimicking the measured A100 controller). Each cell draws $M = 10^4$ null traces; realised rates carry Clopper–Pearson 95% intervals. Figure 6 shows the whole grid.

Level safety holds everywhere. Across the entire campaign no stationary Gaussian, coloured, or non-Gaussian null inflates the realised rate at any stage or level: the failure mode the gate was built to catch — adaptivity silently raising the false-alarm rate — does not occur. On the white null the fixed- α stage is near nominal (0.056, 0.012, and 0.0013 at nominal 0.05, 10^{-2} , and 10^{-3}); on coloured and non-Gaussian nulls the same stage is strongly conservative (realised rate ≈ 0). Tracing the three adaptive steps on the white null, where the fixed- α stage was near nominal, isolates what each one does:

- **Resampling alone (stage 1→2)** deflates the rate by about $3.4\times$ ($0.056 \rightarrow 0.016$ at nominal 0.05). Interpolating onto the angle grid correlates adjacent samples, the Bartlett estimator reads that as extra dependence and inflates $\hat{\Sigma}_\alpha$, and Q_α shrinks. Resampling *colours* the series.
- **Estimating the warp (stage 2→3)** moves the null nowhere ($0.016 \rightarrow 0.016$): the held-out design does what it




Figure 7. The covariance-estimator sweep is the fix-or-reframe axis: realised rate at nominal 10^{-2} versus Bartlett bandwidth b for each shrinkage. The bandwidth that repairs the fully adaptive stage (stage 4) inflates the fixed- α anchor (stage 1). No single configuration calibrates both.

promises.

- **Path selection (stage 3→4)** adds no significant inflation ($0.016 \rightarrow 0.018$; every interval pair overlaps). Even on the resonant AR(2) null, which the Viterbi tracker provably locks onto, the fully adaptive stage fires never in 10^4 traces: locking onto a broad noise resonance does not manufacture the phase-coherent cyclic structure Q_α requires.

The adaptive pipeline is therefore level-safe end to end. What fails is strict two-sided coverage: after resampling the pipeline is uniformly conservative, and the null p -values are not uniform (Kolmogorov–Smirnov uniformity is rejected at every adaptive cell).

No analytic calibration closes the deflation. We swept the open choices — record length, lag set, taper bandwidth, search band, interpolation scheme, and the covariance estimator — selecting on null-calibration quality across the suite, never on power. Only the covariance estimator moves the deflation, and it cannot be closed for the whole ladder at once (Figure 7). Widening the Bartlett bandwidth to $b = 54$ repairs the adaptive stage two-sided (AR(1) realised $0.050/0.012/0.0016$ across the three levels) — confirming that the deflation is an under-bandwidth artefact of the oscillating autocovariance of the resampled series — but the *same* bandwidth breaks the fixed- α anchor, inflating white to 0.073 at nominal 0.05 . No single covariance configuration in the swept family calibrates the ladder end to end, and Kolmogorov–Smirnov uniformity is rejected at every adaptive sweep point regardless. The asymptotic χ^2 tail survives only as a level-safe, power-costing conservative bound; it is not a calibration, and we do not report the adaptive pipeline as an analytic CFAR detector.




Figure 8. Detector bake-off on the cadence-drift axis: detection rate at empirical false-alarm rate versus drift, for each detector and both negative classes. Only the tracked order family retains power against the structural negatives; the spectral-shape detectors detect the controller third of the mix more strongly than the signal.

Surrogate calibration recovers the null. Running the entire pipeline — tracker, resampling, and test — inside a Fourier-phase surrogate null (which preserves the power spectrum and randomises phases, encoding the linear-stationary invariance) recovers exact calibration where the asymptotic tail was $3\text{--}6\times$ conservative. On the stationary Gaussian nulls the realised rate covers nominal across all three pipeline stages — fixed- α , held-out-warp, and fully adaptive (at nominal $0.05/10^{-2}$: white $0.051/0.009$, AR(1) $0.050/0.013$ fixed- α ; white $0.052/0.013$, AR(1) $0.054/0.010$ fully adaptive; the resonant-AR(2) and held-out-warp cells likewise) — and Kolmogorov–Smirnov uniformity is not rejected, the property the analytic p -values never achieved.³ This is what “calibrated” means throughout the paper: an effective score with per-trace surrogate calibration, not an analytic guarantee.

The controller marks the attribution boundary. The surrogate null encodes linear stationarity, so a genuine limit cycle — which is itself cyclostationary — must and does inflate it: on the measured-controller null the surrogate-calibrated rate rises to $0.17\text{--}0.30$ at nominal 0.05 . This is not a defect. The detector correctly reports coherent cyclic structure in a controller trace; calling that structure *training* is the causal-attribution step that needs the nuisance model of Section 3. The controller is the periodic physical confound that Rung 1 alone cannot exclude, and its false-alarm behaviour is exactly the ceiling the claim ladder assigns to this rung. The separation that Rung 1 *can* deliver — power against this confuser class, not merely against inference — is reported in Appendix C.

C. Detector bake-off

This appendix reports the full grid behind Section 5.1, where the headline findings are stated. The detector set spans the natural alternatives: the fixed multitaper harmonic F-test (mt f); a spectral matched filter; the raw Viterbi tracker score; the fixed- α Dandawaté–Giannakis statistic given the *true* f_0 as an oracle; and three variants of the tracked order pipeline — split (held-out warp), full (fully adaptive pooled), and semi-coherent (a per-block Fisher combination). We score positives (Ko training traces) over a cadence-drift axis $f_0^{\text{drift}} \in \{0, 0.1, 0.2, 0.4, 0.8, 1.5\}$ Hz against two negative classes, each with its own empirical false-alarm threshold: the modelled inference null, and a mixed structural population (equal parts AR(1), heavy-tailed AR(1), and the measured controller). With $n = 200$ traces per class the 10^{-2} threshold rests on the second-highest negative score — a granularity we state rather than hide. Figure 8 plots the power curves; Table 4 gives the detection rates at nominal false-alarm rate 10^{-2} .

³Frozen on the MATS compute cluster over stages $1/3/4 \times$ eight nulls at the surrogate sizing $M = 2000$ null traces $\times S = 999$ surrogates (the level-safety and bake-off numbers above are at $M = 10^4$; surrogate cells are capped lower because each carries S inner resamples). This supersedes the earlier partial $M = 1000/S = 199$ grid. The heavy-tailed but *linear* non-Gaussian nulls stay level-safe (conservative: t -innovation AR(1) $0.037/0.005$, log-normal $0.035/0.006$ fully adaptive); the one exception is the *nonlinear* squared-Gaussian null, which the fully adaptive stage over-rejects mildly ($0.087/0.018$, $\approx 1.7\times$, below the pre-registered $3\times$ threshold). Its cause is the same linear-stationarity assumption analysed next for the controller: a Fourier-phase surrogate cannot reproduce the quadratic phase structure of a squared process and so under-spreads the null. The level-safety claim is thus scoped to linear-stationary nulls.

Table 4. Bake-off detection rate at nominal false-alarm rate 10^{-2} , by detector and cadence drift (Hz), against the inference null and the mixed structural population ($n = 200$ per class).

Detector	0	0.1	0.2	0.4	0.8	1.5
<i>vs. inference null</i>						
mtf (fixed multitaper F)	0.02	0.01	0.01	0.01	0.02	0.00
spectral matched filter	0.94	0.53	0.32	0.20	0.31	0.33
Viterbi tracker	1.00	1.00	1.00	1.00	1.00	0.98
DG order (split)	0.28	0.17	0.04	0.01	0.01	0.01
DG order (full)	0.73	0.62	0.55	0.23	0.06	0.01
DG order (semi-coherent)	1.00	1.00	0.97	0.56	0.07	0.01
DG fixed- α (oracle f_0)	0.34	0.07	0.01	0.01	0.00	0.00
<i>vs. mixed structural nulls</i>						
mtf	0.00	0.01	0.00	0.00	0.01	0.00
spectral matched filter	0.00	0.00	0.00	0.00	0.00	0.00
Viterbi tracker	0.00	0.00	0.00	0.00	0.00	0.00
DG order (split)	0.30	0.17	0.04	0.01	0.01	0.01
DG order (full)	0.80	0.70	0.61	0.31	0.10	0.01
DG order (semi-coherent)	1.00	1.00	0.99	0.69	0.14	0.03
DG fixed- α (oracle f_0)	0.36	0.10	0.03	0.01	0.01	0.00

Two details behind the summary in Section 5.1 are worth recording here. The hard case sharpens the confuser finding: at drift 0.4 Hz against controller-only negatives the order pipeline holds 0.8–1.0 detection at 10^{-2} while every spectral-shape detector, including the fixed- α oracle, drops to near zero — a Viterbi score happily locks onto a limit cycle, and only phase coherence at the tracked order distinguishes the two. And among the tracked variants the *semi-coherent* one dominates the pooled statistic at every drift, filling the fast-wander power hole up to about 0.4 Hz while remaining level-safe.

The fast-wander boundary is worth stating precisely. Beyond drift 0.8 Hz every calibratable detector fades (the semi-coherent variant falls to 0.07–0.14 at 10^{-2}); only the raw Viterbi score keeps power against the inference null there, and it cannot tell a training run from a controller. A 300 s record does not resolve fast cadence wander with an honest null. This is a scope statement for the paper, not a repairable defect.

The optimality ceiling. Because we own the generators, we can benchmark against an optimality ceiling: an approximation to the Neyman–Pearson optimal test between the training and null generators. No closed-form likelihood exists — both are black-box samplers — so we use a Whittle spectral likelihood ratio on the in-band periodogram, with a per- f_0 Monte-Carlo training template bank marginalised over cadence, fitted on a disjoint stream and applied to the *exact* bake-off evaluation populations (parity check: maximum detection-rate difference 0).⁴ The ceiling is perfect: area under the curve and detection rate are 1.0 in every column at every drift (Figure 9). The training comb and the line-free nulls are separable in principle from a 300 s trace, so all the difficulty the bake-off exposes lives in the detector, not the problem.

Read as a fraction of this ceiling — the reading summarised in Section 5.1 — it is `dg_order_full` that carries the tracked family against the structural confusers, reaching 0.93 of the ceiling at low drift and fading to 0.10 by the fastest wander, while the Viterbi tracker and matched filter score zero there because they lock onto the limit cycle. The widening gap to the ceiling at high drift is the honest headroom above today’s detectors, not a limit of the channel. The line-tracking literature names the natural candidate for that headroom: under the hidden-Markov wander model the tracker already assumes, the Viterbi score is the best-single-path approximation, and the Neyman–Pearson statistic is the forward-algorithm marginal over *all* frequency paths (Streit & Barrett, 1990); evaluating that sum-over-paths statistic on these axes is future work. The one caveat runs in our favour: the ceiling is optimal only under the Whittle model, a second-order stationary-Gaussian statistic that discards harmonic-phase coherence and non-Gaussian structure, so the true optimum can only exceed it. The perfect ceiling is therefore a lower bound on what is achievable, and reporting deployable detectors as a fraction of it is conservative in the right direction.

⁴Frozen on the MATS compute cluster (slurm job 5646), $n = 200$ per class, 61 cadence templates, 2000 Monte-Carlo draws each; results/st1/np.ceiling.summary.json.




Figure 9. Corpus-free detectors as a fraction of the Whittle Neyman–Pearson ceiling versus cadence drift, against the inference null and the structural confusers. The ceiling is 1.0 everywhere; the Viterbi tracker saturates it against the inference null, while only the tracked order family approaches it against the structural confusers. The gap that opens at high drift is the headroom above current detectors.

D. Full attack sweep

The main text reports the frontier as a single cost-of-hiding overlay (Figure 5). This appendix gives the full sweep: detection against physical de-periodicisation for all seven attack families and both detector classes (Figure 10), the per-detector response to variable real work (Figure 11), and the cost-of-hiding view against each detector class (Figure 12). Here the tracking class is the best of the Viterbi tracker and the two order-tracked Dandawaté–Giannakis statistics of Appendix B; on every cell with a measured cost its envelope equals the Viterbi tracker’s alone.

E. Meter boundary details

This appendix collects the meter-boundary material behind Section 5.2: the full per-detector grid (Figure 13), the transfer-function notch sub-sweep, the decomposition of the frontier’s meter-erasure point, and a comparison with measured board-level work. The robustness ordering across the grid is Viterbi $\geq$ spectral $>$ the DG order family, which is the most sensitive detector at the reference channel but the first to die as the channel degrades; the multitaper F-test is at its floor throughout.

The frontier’s meter-erasure point, decomposed. This sweep also resolves the single integrating-1 Hz point reported on the frontier, where even the tracker degrades (Viterbi detection rate 0.32, every other detector near chance). That point superposes two independent channel defects, and the boundary map shows either one suffices: at the reference channel the tracker is at 1.0; adding the 1 s integration alone drops it to 0.25; the 1 Hz sampling alone drops it to 0.36; the two together give 0.30. The meter-erasure regime is over-determined — it is a statement about the instrument, and a demand on it, rather than a de-periodicisation attack the prover can mount for free.

The specification against measured hardware. Work at the other end of the bandwidth range puts this requirement in perspective. A verifier-owned probe on a single H200’s supply conductors finds the training signature concentrated below about 20 kHz and still classifiable after decimation to about 38 kS/s, where that study stops (Gargiulo & Kulp, 2026) — three to four orders of magnitude above the rates swept here. The two specifications are complementary rather than comparable, because they are requirements on different observables: theirs is the single-card optimizer-update phase, ours the collective iteration cadence. What the measured result does supply is real-hardware support for the direction of the collapse we find, since the fine structure a board-level probe exploits is long gone by the time a facility meter reads, and whatever survives at a few hertz must therefore be a slower and more collective observable than the one a probe sees.

F. Conditional classification against the inference null

The main text detects training against the inference null with a single statistic. This appendix reports a fuller classification task built on the same detectors — an eight-feature vector, three decision rules, six domain shifts, and seven semantic controls — whose control outcomes are summarised in Table 2.

ST2 de-periodicisation frontier (% = measured throughput overhead)




Figure 10. The de-periodicisation frontier: physical de-periodicisation (cadence coefficient of variation) against best tracking and best fixed detector power at nominal false-alarm rate 0.05, over all attack families. The near-vertical drop of the fixed class at the variable-real-work family, at essentially zero throughput cost, while the tracking class holds, is the scope boundary; costlier families (idle jitter, drift, shaping) sit at high throughput overhead without erasing the tracked line.

F.1. From detection to classification

Rung 2 sharpens the structural question into a decision: is a trace more training-like than a stated inference population? A decision needs a reference, and ours is the modelled continuous-batching null of Equation (6); every Rung-2 conclusion is conditional on it. We fix the task in advance — the feature vector, the decision rules, and the reporting — so that the result is not retrofitted to the answer.

Each trace is reduced to an eight-dimensional physics feature vector estimated *from the trace alone*, with no oracle knowledge of the cadence: the six Rung-1 detector scores (multitaper F, harmonic comb, spectral matched filter, Viterbi, fixed- α DG, and tracked-order DG), plus two stability features read off the estimated Viterbi path — its frequency-path stability and its phase-folded waveform repeatability. On this fixed vector we compare three decision rules of increasing freedom: (i) a prespecified *physics* score, a fixed linear combination with the fixed-cadence comparators pre-registered out; (ii) a logistic discriminant *fitted* to the scenario generator and scored out-of-fold; and (iii) a random-forest *learned* reference, included to expose how much separation comes from the chosen null rather than from the structural method. We report area under the ROC curve and the false-negative rate at a false-positive rate of 0.05, kept separate throughout from the Rung-1 structural false alarms.

The population the rules are evaluated on is fixed in advance too. They are fitted on the nominal setting and evaluated zero-shot on six domain shifts, each moving one factor: device headroom down to 0.70 and up to 0.95 of peak; observation duration halved to 150 s, which halves the frequency resolution; the cadence band shifted up to 0.8–1.8 Hz, so that its top now lies outside the 0.3–1.7 Hz search band; a meter carrying the controller limit cycle; and a meter with heavy coloured noise and baseline wander.

F.2. Classification against the inference null

We now apply the three decision rules of Section F.1 to the stated population, the six domain shifts, and a set of semantic controls whose training status we know.⁵

⁵Frozen on the MATS compute cluster (slurm array 5612), $n = 200$ per class, 300 s traces, band 0.3–1.7 Hz, 14 cells (the stated population, six transfer shifts, and seven semantic controls); `results/rung2/rung2-summary.json`.




Figure 11. Detector power versus variable-real-work level for each detector. The fixed matched filter and multitaper F-test cross below chance by level 0.35; the tracked order statistics degrade more slowly; the Viterbi tracker holds full power until it finally bends at level 0.7 (detection rate 0.59) — a level whose throughput and learning cost we do not measure.

Stated population and transfer. On the stated population all three rules separate training from the hard inference null perfectly: area under the curve 1.00 with zero false negatives at a false-positive rate of 0.05 (Figure 15). The separation is not brittle to the population: across all six domain shifts of Section F.1, evaluated zero-shot, the area under the curve stays 1.00. The physical schedule of efficient training is far enough from the inference null that leaving the stated parameter box does not close the gap.

The semantic controls set the ceiling. The controls are the point of the rung. We built seven loads whose training status we know and asked what each rule scores (Table 2). Five non-training loads — a training-shaped non-ML kernel loop, a controller limit cycle, a periodic inference request generator, a gradient-only computation, and a discarded-update decoy that runs the full forward/backward/optimizer step and throws the update away — *all* score as training, at fraction 1.00 under every rule. A genuine training run diluted by an unrelated co-resident load is still caught (1.00). The reading is the one the claim ladder assigns to this rung: from power alone the verifier can certify a *training-shaped periodic compute schedule*, not that the computation is gradient-descent training rather than a mimic. The first three loads are the practically binding ones, because they need no adversary: the controller cycle scoring as training is the same attribution boundary Rung 1 met in Section B.3, and calling structure training needs the nuisance model of Section 3. The decoy is the extremal member of the same class and the one that cannot be engineered away — Section 7.1 rules out any passive detector that separates it — which is why this ceiling motivates the active protocol of Appendix H rather than a better classifier. The ceiling is a property of the observable rather than of the meter’s bandwidth. A learned classifier reading a board-level probe at 10 MHz separates

Cost of hiding: what measured adversary cost buys against each detector class




Figure 12. Cost of hiding: measured throughput overhead against hiding ($1 - \text{detection rate at false-alarm rate } 0.05$) for the four cost-anchored attack families, with the Pareto staircase — the best hiding available at or below a given cost. Against the tracking class (left) the staircase stays on the floor across the whole measured cost axis; against the fixed class (right) the variable-real-work family reaches hiding ~ 0.74 at zero cost. Cells whose attack level has no hardware cost anchor cannot be placed on the cost axis and are shown in the hatched strip, among them the level at which the tracker itself begins to fail. Attack levels are labelled where they hide something; the rest sit on the floor.

training from inference by the spectral signature of the memory-bound optimizer-update phase (Gargiulo & Kulp, 2026); the discarded-update decoy runs exactly that phase and discards only its result, so the mechanism such a classifier keys on is present in a load that is not training. More bandwidth resolves the physical schedule more sharply, and it is the schedule, not the semantics, that a passive meter reads.

Asynchronous training: a robustness result, not a miss. One control, strongly de-periodicised genuine training, was pre-registered as an expected miss — the case scoped out of Rung 2. At $n = 200$ it is not missed. The tracked physics rule catches it (area under the curve 0.89, 72% flagged) and the learned reference catches all of it (100%); only the fitted linear discriminant is weak (0.75, 55%). This is consistent with the tracker-robustness story of Rungs 1 and 2: the asynchronous control de-periodicises by drifting the cadence centre and slipping the phase, which leaves a wandering but still coherent line — exactly the regime order tracking was built to survive. We therefore report it as a robustness result rather than the boundary. The genuine evasion is not slow wander but the $\approx$ zero-cost variable-real-work attack of Section 6, which moves the cadence *without* leaving a coherent line, and which is where the tracker actually bends.

G. Identifiability details

This appendix states the cost-of-hiding bound summarised in Section 6, proves it and the non-identifiability result of Section 7.1, and reports the numerical check that fixes their constants.

G.1. The cost of hiding

Two questions have to be kept apart. *Detection* asks whether this trace is training rather than the null, and its tools are likelihood-ratio and cyclostationary tests. *Cost of hiding* asks how far the prover must distort its schedule before the training-trace *distribution* becomes indistinguishable from the null *distribution*, and its tools are distributional distances. The distinction matters because our empirical results bound the wrong direction for a governance claim: a detection sweep says how much can be hidden from *our* detector, which a better detector moves. This subsection states what can be argued




Figure 13. Minimum meter specification: detector power at nominal false-alarm rate 0.05 over sample rate $\times$ integration window, one panel per detector, honest training against the hard inference null. The live region is sample rate ≥ 2 Hz and integration window ≤ 0.5 s; below either edge the tracking class collapses toward chance.

Notch sub-sweep (20 Hz sampler): detection vs in-band anti-resonance




Figure 14. Transfer-function notch sub-sweep at 20 Hz sampling: tracker power versus notch centre frequency and depth. Only a deep notch placed at the cadence centre (1.0–1.2 Hz) materially degrades detection, and it halves rather than erases it.

in the other direction. It is the one part of the paper that does not come from the generator; the derivations behind it follow in Section G.2.

What can be argued quantitatively concerns the *cost* of hiding, and it is stated over the population this paper scopes itself to: efficient, synchronous, iteration-structured training whose compute/communication alternation exposes a recoverable cadence f_0 , and in which the synchronous exchange occupies a down-phase whose duration does not scale with the work accumulated in the preceding up-phase. Call an attack ε -covert against a class of verifiers if no test in the class attains detection advantage — true-positive minus false-positive rate — above ε . Any single test’s advantage is at most the total variation $\text{TV}(P_{\text{train}}, P_{\text{null}})$ between the two laws (Tsybakov, 2009), so covertness against *every* verifier requires that total variation to be at most ε .

Consider the concrete attack the frontier prices: i.i.d. fractional period jitter of scale σ , so iteration i has period $1/(f_0(1 + \xi_i))$ with $\xi_i \sim \mathcal{N}(0, \sigma^2)$. Such jitter *accumulates* — the iteration-boundary phase performs a random walk with diffusion coefficient $D \approx (2\pi)^2 \sigma^2 f_0$ — so the spectral line broadens into a Lorentzian and the coherent power retained in one frequency bin of a record of length T falls off as

$$\frac{P_c(\sigma)}{P_c(0)} \approx \frac{1}{1 + \kappa \sigma^2}, \quad \kappa = \pi^2 f_0 T, \quad (14)$$

a parameter-free prediction that the generator reproduces to within several percent (Appendix G). Now let $J(\sigma)$ be the advantage attained by a fixed-bin threshold test at the nominal cadence, which decreases in σ through Equation (14).




Figure 15. Rung-2 classification on the stated population (ROC / operating points for the three decision rules) and across the transfer shifts. All three rules reach area under the curve 1.00 on the stated population and hold it under every domain shift.

Because $J \leq \text{TV}$, any σ with $J(\sigma) > \varepsilon$ is not ε -covert, so ε -covertiness requires a distortion at least $\sigma^*(\varepsilon)$, the solution of $J(\sigma^*) = \varepsilon$. It is strictly positive and grows as ε falls; measured through the full generator and meter it is $\sigma^*(0.5) = 0.18$ and $\sigma^*(0.2) = 0.36$, which the anchors of Section 6 price at tens of percent of throughput and at order 100% respectively.

Running the argument through an *explicit* verifier is essential, and it is worth saying why, because the familiar covert-communication route goes the other way. Pinsker’s inequality $\text{TV} \leq \sqrt{\text{KL}/2}$ (Tsybakov, 2009; Cover & Thomas, 2006) bounds total variation from above, so forcing that bound below ε would establish that some distortion *suffices* for covertness, not that any is *necessary*. Only a verifier that demonstrably attains advantage J certifies that smaller distortions fail.

Distortion is not yet cost, and the bridge is a physical property of the admissible family that we state as an *assumption* rather than a result, because two steps in its justification are not closed. The *cost-of-hiding assumption* is that within that family, and at fixed total useful work, reaching $\sigma^*(\varepsilon)$ costs throughput or learning efficiency and cannot be free in both. The reasoning is that period jitter then has only two realisations. Inserting idle time inflates the mean period and pays in throughput. Varying the real work per iteration avoids idling, but the down-phase duration does not scale with the accumulated work and that work is quantised, so a component at the exchange cadence survives; removing it requires either idling or desynchronising the exchange, and the latter changes the optimisation being performed.

Three limits fix how far all of this travels, and they are the reason the assumption is labelled as one. First, the surviving exchange-cadence component is recovered by the order-tracking class, not by the fixed-bin verifier that supplies σ^* : under variable real work Section 6 measures the fixed tests below chance. Read strictly, the argument shows that a work-varying prover stays visible to a *tracker*, not that it must pay to become ε -covert against the verifier whose σ^* we computed. What the numbers do support is the comparative claim — at matched total distortion, work variation is the more detectable of the two at every $\sigma \geq 0.05$, the two being indistinguishable within sampling error below that (Figure 17). Second, fixing total useful work is a real restriction: substituting useful non-training work, the additive-burial and dilution families above, modulates the period while idling nothing and degrading no optimisation. The disjunction is therefore asserted over timing




Figure 16. Semantic controls: fraction scored as training per decision rule. The non-training decoys are indistinguishable from training under every rule (the semantic-decoy ceiling); only asynchronous genuine training separates the rules, and the tracked physics rule still catches most of it.

attacks at fixed work. Third, σ^* is read off one exhibited verifier and J is measured between the two laws of our own generator. Because $J \leq \text{TV}$, a σ ruled out by the fixed-bin test is ruled out against any verifier class containing that test — including the five-detector bank of Section 6, order tracker and all — but a stronger verifier attains a larger advantage at each σ and so yields a larger σ^* . Since our own experiments show the tracking class to be the stronger one, the threshold we report is a loose reading of the true requirement. And the claim throughout is that hiding *to a stated* ε must be paid for, expressly not that de-periodicisation is inefficient in general, which Section 6 refutes by exhibiting a near-free attack that defeats every fixed test.

G.2. Derivations and numerical check

Why no passive detector identifies semantics. A passive verifier observes only P_{obs} , so its decision is a measurable function ϕ of that trace and its rejection probability is $\mathbb{E}_{K_h(w)}[\phi]$. If $K_{h_0}(w_0) = K_{h_1}(w_1)$ the two expectations coincide for every ϕ , so no test separates the pair, and no post-processing of the trace can (data processing cannot create distinguishability that the laws do not contain). The claim about restrictions follows by contraposition: a non-vacuous guarantee must exclude the equivalent pairs, and each listed lever excludes some of them. The amplitude degeneracy of Equation (10) shows the pairs are not exotic. The meter constrains only the product $r(t)F(t)$, so a workload with a different operation rate and a compensating energy cost per operation is already observationally equivalent, and temporal structure — not amplitude — is what the trace identifies. The near-constructive same-hardware case is the discarded-update decoy: forward, backward, and optimiser kernels executed on the declared schedule with each updated state discarded reproduces the compute/communication alternation, hence the trace, without retaining a model. The converse direction is equally open to the prover: genuine training can be run asynchronously, diluted among other loads, or actively power-shaped, so the

equivalence classes are populated from both sides.

Line attenuation under phase diffusion. Boundary i nominally carries phase $2\pi i$; fractional period jitter perturbs each boundary relative to its predecessor, so the *residual* of the boundary phase against a uniform clock accumulates independent increments and performs a random walk with $\text{Var}[\theta(t)] \approx Dt$ and $D \approx (2\pi)^2 \sigma^2 f_0$. For a Brownian phase the periodic component’s autocorrelation decays as $\exp(-D|\tau|/2)$, so its spectrum is a Lorentzian rather than a line, with half-width at half-maximum $\gamma = D/4\pi = \pi\sigma^2 f_0$ in Hz. A record of length T resolves frequency only to $1/T$, so the coherent power a fixed-bin verifier collects is the Lorentzian mass within a half-width $B = 1/2T$ of the centre,

$$\frac{P_c(\sigma)}{P_c(0)} = \frac{2}{\pi} \arctan\left(\frac{B}{\gamma}\right) = \frac{2}{\pi} \arctan\left(\frac{1}{2\pi f_0 T \sigma^2}\right), \quad (15)$$

whose small-argument form is $1/(\pi^2 f_0 T \sigma^2)$. Matching that tail to $(1 + \kappa\sigma^2)^{-1}$ gives

$$\kappa = \pi^2 f_0 T, \quad (16)$$

a parameter-free prediction rather than a fitted scale — the width and bandwidth conventions are fixed by the explicit calculation for the rectangular-bin, untapered estimator used below, so no $O(1)$ factor is left free. The Lorentzian of Equation (14) is not a pointwise approximation to Equation (15): the two differ by up to 22% near $\sigma = 0.02$, where the arctan is still turning over, and agree to better than 1% only for $\sigma \gtrsim 0.2$. What matters is that using the Lorentzian as the fit model does not bias the constant it returns — fitting Equation (14) to Equation (15) across the σ grid swept below recovers $\kappa = 2955$, within 0.2% of $\pi^2 f_0 T$.

The distortion covertness requires. Let S be the fixed-bin statistic and consider the threshold tests $\{S \geq c\}$. The advantage of the best such test is $J = \sup_c (\Pr_{\text{train}}[S \geq c] - \Pr_{\text{null}}[S \geq c])$, the *one-sided* Kolmogorov–Smirnov distance of the score laws (one-sided because training traces score high, so only that direction is a detection advantage; one-sided $\leq$ two-sided $\leq$ TV, so the inequality we need still runs the right way). These are probabilities under the two laws, so J is a population quantity, and each $\{S \geq c\}$ is a deterministic single-trace test that attains its own difference — existence is all the definition of ε -covert needs. Each such test is a measurable function of the trace, so by data processing $J \leq \text{TV}(P_{\text{train}}, P_{\text{null}})$. If $J(\sigma) > \varepsilon$ then total variation exceeds ε and the attack is not ε -covert; equivalently, ε -covertiness requires $J(\sigma) \leq \varepsilon$. Empirically J decreases in σ over the swept range, so the requirement becomes a lower bound $\sigma \geq \sigma^*(\varepsilon)$; we do not prove that monotonicity, which would need a stochastic-ordering argument for the score laws rather than merely the decay of the coherent power that drives them. Fitting the functional form of Equation (14) to J , $J(\sigma) = J_0/(1 + \kappa_J \sigma^2)$, and inverting gives

$$\sigma^*(\varepsilon) = \kappa_J^{-1/2} \sqrt{J_0/\varepsilon - 1}, \quad (17)$$

which grows like $\varepsilon^{-1/2}$ for small ε . That asymptotic form is a property of the *fitted* rolloff extrapolated beyond the swept range, not a derived law, which is why Section G.1 claims only that σ^* is positive and increases as ε falls. Enlarging the verifier class can only increase the attainable advantage at each σ , hence only raise the class-indexed $\sigma_{\mathcal{C}}^*(\varepsilon) = \inf\{\sigma : \sup_{\phi \in \mathcal{C}} J_{\phi}(\sigma) \leq \varepsilon\}$; σ^* above is that quantity for the singleton fixed-bin class.

Two points of bookkeeping. First, the rolloff constant κ_J of the *advantage* is not the κ of the *coherent power* in Equation (14): the advantage saturates near one while any excess at f_0 remains separable from a line-free null, so it decays far more slowly. The two are not interchangeable in either direction — coherent power is not a detection advantage at all, so it can never be substituted into Equation (17), whatever its rolloff. Second, J is not interchangeable with the rank-based Gini index $2 \text{AUC} - 1$: the latter can exceed J , so it is not the advantage of any single-trace test and cannot appear in Equation (17). Both are reported in the summary artefact; only J carries the bound.

The cost-of-hiding assumption, and why it is an assumption. Hold the prover’s total useful work fixed. Then a period is occupied either by training work or by idling, and period jitter has two realisations. Padding with idle time holds the real work per iteration fixed, so every period must be at least as long as the compute it contains and the mean period inflates with σ ; throughput falls, and Section 6 measures the penalty at 15–375% over the swept range. Varying the real work per iteration avoids idling and is measured at essentially zero throughput cost, but within the admissible family it cannot make the schedule aperiodic: the accumulated work is an integer number of micro-steps, so the up-phase is quantised rather than continuously tunable, and the exchange down-phase keeps its duration whatever that count, leaving a component at the

exchange cadence. Removing *that* requires either inserting idle time, which returns to the first realisation and its throughput cost, or desynchronising the exchange, which alters the optimisation actually performed and is therefore paid for in learning efficiency.

Two gaps stop this being a proof, and both matter.

The first is that the surviving component is not one the fixed-bin verifier can use. Section 6 measures the fixed tests under variable real work falling to detection rate 0.26 with area under the curve 0.44 — below chance — and the multitaper F-test at its floor. So the residual at the exchange cadence is recovered only by the order-tracking class, which is precisely the class the fixed-bin argument does *not* instantiate. Read strictly, then, the argument shows that a work-varying prover remains detectable *by a tracker*, not that it must pay to become ε -covert against the fixed-bin verifier whose σ^* we computed. Closing that gap needs σ^* for the tracking class, which we do not have.

The second is that fixing total useful work is a real restriction. A prover free to substitute useful *non-training* work — the additive-burial and dilution families of Section 3.2 — can modulate the period while idling neither the device nor the training job, and pays in neither training throughput nor learning efficiency. Frequency scaling is a further route the two-realisation dichotomy does not cover. So the disjunction holds over timing attacks at fixed work, not over the adversary’s full repertoire.

What the numerical check does support is the weaker comparative claim: at matched total distortion, and for every $\sigma \geq 0.05$, the work-varying attack is *more* detectable than idle-padded jitter, so substituting real-work variation does not buy coventness for free relative to the priced attack — it trades throughput cost for detectability. That comparison is measured; the cost disjunction itself is assumed, and labelled accordingly.

Numerical check. Figure 17 tests every load-bearing step against the generator and meter of Section 2 rather than an idealised channel. Training and inference-null populations ($n = 600$ per class, $f_0 = 1$ Hz, $T = 300$ s, $f_s = 20$ Hz, seeded) are pushed through the full observation map and scored with the fixed-bin statistic. The coherent power follows Equation (14) with $R^2 = 0.95$ and a fitted $\kappa = 3068$ against the parameter-free $\pi^2 f_0 T = 2961$ of Equation (16), with nothing fitted — but that 4% should be read as an order-of-magnitude agreement rather than a two-figure one. The linearised fit weights each point by σ^4 , so 81% of its leverage sits on the largest σ swept, and there the generator realises a cadence CV of 1.9 against a nominal 0.5: the mean period has drifted well off f_0 and the fixed bin is no longer measuring line broadening alone. Restricting the fit to the faithful regime $\sigma \leq 0.1$, where the realised CV tracks the nominal to within about 13%, gives $\kappa = 2812$, or 5% below the prediction. The honest statement is agreement at the several-percent level from either end of the sweep, which still separates $\pi^2 f_0 T$ from $\pi f_0 T$ by the factor of 3.1 that distinguishes them — the strongest single confirmation of the physical picture here.

The attained advantage requires more care, because it is estimated. Our plug-in $\hat{J}$ maximises over thresholds on the same sample it evaluates, so it is upward-biased; under H_0 at $n = 600$ its fluctuation reaches 0.093, and a measured $\hat{J}$ below that carries no evidence of separation. Because the fit linearises as $J_0/J - 1$, which diverges as $J \rightarrow 0$, such points would otherwise dominate it — at the $n = 80$ sizing we first used, two washed-out points carried 97% of the fit and κ_J varied fourfold across seeds. We therefore exclude points below the H_0 floor and report intervals. On the surviving window $\sigma \in [0.05, 0.35]$ the advantage rolls off with $\kappa_J = 32$ (95% bootstrap interval [24, 40], $R^2 = 0.90$, $J_0 = 1.0$), giving

$$\sigma^*(0.5) = 0.18 [0.16, 0.20], \quad \sigma^*(0.2) = 0.36 [0.32, 0.41],$$

and, by extrapolation beyond the swept range, $\sigma^*(0.1) = 0.53$ and $\sigma^*(0.05) = 0.78$.

Two caveats bound how far these numbers travel. First, the generator realises the nominal distortion only up to $\sigma \approx 0.2$: beyond it the reciprocal Gaussian period draw and a clamp take over, and the measured cadence coefficient of variation runs high (1.2 at $\sigma = 0.35$, 1.9 at $\sigma = 0.5$), as does D against $(2\pi)^2 \sigma^2 f_0$. Only $\sigma^*(0.5)$ falls inside that faithful region; $\sigma^*(0.2)$ and everything beyond it do not, and the artefact flags them. Second, the fit window reaches past $\sigma = 0.2$ by choice rather than necessity: the five points at $\sigma \leq 0.2$ are unsaturated and do support a fit, but the deep-rolloff point at $\sigma = 0.35$ carries most of the leverage on κ_J . Restricting the fit to the faithful region alone gives $\kappa_J = 22$ and hence larger thresholds, $\sigma^*(0.5) = 0.22$ and $\sigma^*(0.2) = 0.43$. The thresholds are therefore good to about 20%, not to two significant figures, and the direction of that fit-window uncertainty is towards *more* required distortion, not less. Finally, the work-varying series attains a higher advantage than period jitter at every $\sigma \geq 0.05$, the two being indistinguishable within sampling error below that, which is the comparison the cost-of-hiding assumption appeals to. Both arms set the period-jitter scale explicitly, so this is a comparison at matched total distortion.

Lower-bound spike: covertness of i.i.d. phase jitter (fixed-bin verifier, $f_0=1$ Hz, $T=300$ s)




Figure 17. Numerical check on the covertness cost bound, through the full generator and meter ($n = 600$ per class). **Left:** coherent power at the nominal cadence against jitter scale. The fitted and the parameter-free analytic rolloffs of Equations (14) and (16) are indistinguishable at this scale; the fitted constant agrees with the parameter-free prediction at the ten-percent level (see the numerical check). **Right:** the advantage attained by the fixed-bin threshold test, $J \leq TV$, for idle-padded period jitter and for the near-free work-varying attack; dotted lines mark $\varepsilon \in \{0.1, 0.5\}$, whose crossings give $\sigma^*(\varepsilon)$ (the $\varepsilon = 0.1$ crossing lies beyond the swept range and is reached by extrapolation). The dash-dotted line is the H_0 fluctuation of the estimator: the $\sigma = 0.5$ point falls below it and is excluded from the advantage fit shown here (it is retained in the coherent-power fit on the left, which has no sampling floor of its own). For $\sigma \geq 0.05$ the work-varying attack is the more detectable of the two at matched distortion, the comparison the cost-of-hiding assumption appeals to; below that the two are indistinguishable within sampling error. Both arms set the period-jitter scale explicitly, so the distortion really is matched. The verifier here is fixed at the known cadence, not the order tracker of Section 4.

From distortion to a price. The distortion thresholds can be read against the measured cost anchors of Section 6, which price the idle-padded realisation on the same nominal σ grid. Hiding from the fixed-bin verifier to $\varepsilon = 0.5$ needs $\sigma^* = 0.18$, which falls between the measured levels $\sigma = 0.1$ and $\sigma = 0.2$ and so between 32% and 66% throughput overhead; to $\varepsilon = 0.2$ it needs $\sigma^* = 0.36$, between 159% and 375%. Given the roughly 20% uncertainty on σ^* established above, and the scope limits set out in Section G.1, the robust form of the statement is the order of magnitude rather than the bracket: driving a fixed-bin verifier to $\varepsilon = 0.5$ costs *tens of percent* of throughput, and to $\varepsilon = 0.2$ costs *of order 100% or more*. Both readings survive the fit-window sensitivity; the individual brackets do not, and the $\varepsilon = 0.5$ one in particular moves up a level under the faithful-region fit.

This is the sense in which the bound has a price, and Figure 12 is its empirical counterpart. Three things keep the two from being conflated. The thresholds come from a fixed-bin verifier at the known cadence on 600 traces per class, whereas the frontier scores a five-detector bank on 200. The pairing is by nominal σ , but the two campaigns realise the same nominal setting somewhat differently once past the faithful region — cadence coefficient of variation 0.50 here against 0.84 there at $\sigma = 0.35$ — so the mapping is exact only for $\sigma \lesssim 0.2$. And the frontier prices attacks against a detector bank that includes the tracker, whose σ^* we never computed. They are read side by side, not on shared axes.

What the bound does not cover. The propositions above concern one attack family. Making the measured frontier a statement about hiding in general would need one of two further arguments: that the swept families span what a hiding adversary can do, or a bound of the kind above for each of them. We have the second for i.i.d. period jitter and make no claim of the first, so Section 6 should be read as a sample of the attack space rather than its boundary. The choice of distributional distance is also not settled. We use total variation, reached through an attained advantage; Kullback–Leibler and Hellinger distances tensorise more conveniently over the many-iteration product law and may give a tighter closed form, which we leave open.

H. From passive evidence to active authentication (Rungs 3–4)

Rung 2 certifies the physical schedule, not training semantics. Its ceiling is semantic non-identifiability: loads that reproduce the compute–communication cadence of training without training score as training under every passive rule (Section F.2), and no passive test separates them, because the meter sees only the schedule and the mimic copies the schedule. The discarded-update decoy is the sharpest instance — Section 7.1 rules out any passive detector that separates it — even though a concealing prover would never mount one (Section 3.2). The way past a passive ceiling is to stop watching and start asking: authenticate a physical device by its response to an unpredictable query, the premise of distance-bounding protocols (Brands & Chaum, 1994), physical unclonable functions (Pappu et al., 2002) and, in compute governance, of challenge–response telemetry (Monfared et al., 2026). If the verifier perturbs the machine and the machine must react, a mimic that merely replays a recording cannot follow: it does not know what it was asked.

Rung 3 makes this concrete with a challenge-synchronous protocol. The verifier emits an unpredictable balanced challenge sequence $c_k \in \{-1, +1\}$; the prover runs a committed, semantically-equivalent variation of its training schedule keyed to c_k — for instance a small pre-agreed shift in the boundary between two iterations that leaves the computation performed unchanged. The verifier correlates the external power trace against c_k over a causal lag window, in the manner of a lock-in amplifier or a spread-spectrum receiver recovering a known code from noise. The null needs no corpus and no matched hardware: it comes from re-randomising the challenge and recomputing the statistic. A positive result authenticates a physical fact — a load behind the meter responded to a signal it could not have predicted. This escapes the semantic decoy, whose passive schedule is by construction uncorrelated with the fresh challenge (Section 3.2). Its own ceiling is the challenge-aware dummy load: a device that receives the challenge and mimics the prescribed response without training. Unlike the passive decoy, this one is an attack a prover has reason to mount, because Rung 3 asks it to demonstrate training rather than to hide it.

The correlation test is exact under the sharp null that the challenge has no effect on the measured process — but only conditional on the randomisation design, and the guarantee is fragile to the usual abuses (Fisher, 1935). Any maximisation the verifier performs, over causal lags, candidate impulse responses, or preprocessing choices, must be repeated inside every re-randomised replicate, or the reported level is optimistic. Carryover between challenge periods, any adaptivity in the prover’s schedule choices, and balance constraints on the challenge sequence all belong in the design that the re-randomisation samples from. We also flag the weaker reading a careless protocol invites: over a long enough record a correlation detector accrues the $\sqrt{N}$ processing gain of any matched receiver, so a positive result must be shown to certify the challenge binding and not merely that the verifier integrated for long enough.

Rung 3 authenticates a physical response; it does not tie that response to training. Rung 4 adds that binding, and it is the rung this paper does not build. Rather than gesture at “a proof of training,” we state the missing piece as an ideal functionality (Canetti, 2001): a sound, fresh, challenge-bound transcript or proof that ties each challenge bit to the prescribed semantically-equivalent variation of declared training state transitions. Such a primitive must provide commitment before the challenge is revealed, unpredictability at commitment time, sound binding of the response to valid declared state transitions, freshness against replay, and an explicit position on whether the responding load and the declared computation are co-located. Proof-of-learning constructions (Jia et al., 2021) are the closest existing candidate, but they assume access to the training trace rather than an external meter. Conditional on such a primitive, challenge-synchronous power supplies physical evidence for the declared response; without it, the semantic work binding remains an assumed, open layer.

Two ceilings therefore remain. Rung 3’s is the challenge-aware dummy load; Rung 4’s is co-location: a genuine remote proof combined with a local challenge-correlated dummy still satisfies both the correlation test and the transcript unless the assumed primitive explicitly rules the separation out. Closing them empirically is future work, and it needs no cluster — challenge modulation is legible on a single card. A full Rung-3 study would fix a concrete pair of semantically-equivalent schedules, run the complete search inside every re-randomised replicate, pit the detector against a challenge-aware dummy-load adversary, and trace detection against modulation depth, throughput overhead, and observation time. We leave it as the active-verification direction the passive results here motivate but cannot reach — the response to the ceiling that Section 7.1 makes unavoidable for any passive detector (Section 7).

I. Scratch

Material cut from the body during the 2026-09 restructure and parked here pending a decision. **This appendix is not part of the paper.**

I.1. Contributions list (cut from the introduction)

The introduction previously closed with an enumerated contributions list. It was cut because the section it headed is now prose; the claims themselves are made and evidenced in the body.

1. **The claim ladder** (Table 3) as an organising frame for passive compute verification: four claims, their information costs, and their ceilings, closed out one by one in Table 6.
2. **A tracked cyclostationary detector for Rung 1** (Section 4), with a negative methodological result we report rather than bury: no covariance configuration calibrates the complete tracker–resample–test pipeline end to end, so it is not a constant false-alarm-rate test and we do not present it as one. It is an effective score with per-trace surrogate calibration, level-exact on the linear-stationary nulls we tested. A pre-registered bake-off (Appendix C) makes the choice of detector an outcome, and against a Whittle likelihood-ratio ceiling the tracker is essentially optimal on the inference null while only the cyclostationary order family approaches the ceiling on structural confusers.
3. **The de-periodicisation frontier** (Section 6), the central result, priced against measured throughput anchors. Its shape is an asymmetry: against fixed-frequency tests a prover reaches a hiding rate of 0.74 at essentially zero throughput cost by varying real work per iteration, whereas against the tracking class the best anchored attack reaches only 0.16 at 159% overhead.
4. **A minimum meter specification** (Section 5.2): the observation channel is a first-class condition on every rung, and we locate where each detector class dies in sample rate, integration window, and in-band transfer function. A 1 s-integrating, 1 Hz-reporting meter defeats every detector on an unmodified honest workload.
5. **Rung-2 classification with semantic falsification controls** (Section F.2). Three decision rules separate training from the stated inference null perfectly and survive six domain shifts; but five non-training loads — three of them benign operational loads, alongside a decoy that computes gradients and discards them — all score as training. The meter certifies the physical schedule, not training semantics: a ceiling no better passive detector removes, and one that limits positive evidence as much as it admits counterfeits.
6. **A negative transport case** (Section 7.2): what we measured on real hardware, why it does not carry the distributed claim, and what would falsify the prediction that it should.

Supporting these, Section 7.1 gives the formal counterpart of the Rung-1 and Rung-2 ceilings: an observational-equivalence argument for why no passive test can identify training semantics, and a covertness cost bound showing that a prover using i.i.d. phase jitter to reach a target covertness level against an exhibited verifier must pay in throughput or learning efficiency. We are explicit about what that bound does not cover, and Section 6 supplies the near-free attack that shows why the gap matters.

J. Scratch

[Parked] Old Detection-methods introduction. The weakest question on the ladder is whether the trace carries order-coherent power variation that a stated structural null does not produce. Answering it establishes *coherent repeated power structure*, not training: reading that structure as repeated multi-phase computation additionally requires the nuisance model of Section 3 that excludes periodic physical confounds, of which the measured A100 power-controller oscillation (Section 7.2) is the standing counterexample. This section develops the detectors and the calibration machinery; Section 5 reports what they do.

Two prover settings drive the design. A benign prover’s cadence is stationary, and a stationary line has a classical test. An adaptive prover lets the cadence wander, which breaks that test cheaply and forces a different construction. The benign case is easy and the paper is not about the benign case; the adaptive case is the methodological heart, and the burden here is to give its detector a defensible null.

[Parked] Old main-text detector-power subsection.

[Old heading] Detector power and how far it sits below optimal Which detector to use is an *outcome* here, not an assumption. Before scoring any signal we fixed the detector set, the conditions, and the metric, and ran a pre-registered comparison across a cadence-drift axis $f_0^{\text{drift}} \in \{0, \dots, 1.5\}$ Hz against two negative classes: the modelled inference null,

and a mixed structural population of AR(1), heavy-tailed AR(1), and the measured controller. Appendix C reports the full grid; three findings from it carry into everything below.

Tracking is necessary. The fixed- α Dandawaté–Giannakis statistic, handed the *true* f_0 as an oracle, collapses from detection rate 0.34 at zero drift to ≤ 0.1 under any wander — even the true frequency is useless to a coherent fixed- α test once the line moves — and the fixed multitaper F-test has essentially no usable power against either negative class at empirical thresholds. Only the tracked order family separates the wandering line from structural confounds: against the mixed structural negatives the matched filter and the raw Viterbi score have detection rate zero at every drift, because the controller third of the mixture fires them harder than the signal does. Phase-coherent cyclic structure at the tracked order is the one signature the confuser cannot fake, and this is the attribution boundary of Rung 1 drawn with power rather than with false-alarm rate. The honest limit is fast wander: beyond drift 0.8 Hz every calibratable detector fades, and a 300 s record simply does not resolve fast cadence wander against an honest null.

That ranks the detectors against each other but says nothing about how good the winner is. Because we own the generators we can also benchmark against an optimality ceiling. No closed-form likelihood exists — both classes are black-box samplers — so we use a Whittle spectral likelihood ratio on the in-band periodogram with a Monte-Carlo training template bank marginalised over cadence, fitted on a disjoint stream and applied to the exact evaluation populations. The ceiling is perfect: area under the curve and detection rate are 1.0 in every column at every drift. The training comb and the line-free nulls are separable in principle from a 300 s trace, so all the difficulty the comparison exposes lives in the detector, not in the problem.

Read as a fraction of that ceiling, the picture sharpens in a way the frontier of Section 6 depends on. Against the inference null the Viterbi tracker is essentially Neyman–Pearson optimal, holding the ceiling’s power across the whole wander axis, while the fixed matched filter’s fraction of attainable power collapses from 1.0 at zero drift through 0.78 and 0.38 to below zero by drift 0.4 Hz. Against the structural confusers the tracker is *not* enough: only the tracked order family approaches the ceiling, reaching 0.93 of it at low drift and fading to 0.10 at the fastest wander. The widening gap at high drift is honest headroom above today’s detectors rather than a limit of the channel, and the one caveat runs in our favour — the Whittle ceiling discards harmonic-phase coherence and non-Gaussian structure, so the true optimum can only exceed it.

[Parked] Old Discussion opening, ladder lead-in, ladder tables, and governance readings. We set out to ask what a passive power meter can certify about AI training, and organised the answer as a claim ladder in which each rung answers a more governance-relevant question at a higher cost in verifier information, and each rung’s ceiling motivates the next. Table 6 closes the ladder out with the paper’s evidence attached to each rung.

This section fixes what the verifier sees, who is trying to hide from it, and what can be argued about the channel before any detector is run. It also fixes the vocabulary for the claims a meter can support. We organise those claims as a *claim ladder* (Table 3): four claims of increasing governance relevance, each bought at a stated cost in verifier information, and each with a hard ceiling that motivates the next. Rung 1 establishes that order-coherent power structure is present, which is evidence rather than identification. Rung 2 converts that structure into training-likeness at the price of committing to an inference null; its ceiling is semantic non-identifiability. Rungs 3 and 4 leave the passive channel for an active protocol and a transcript primitive, and are specified but not run (Appendix H). The ladder is a ladder of claims and assumptions, not a quantitative bound on unidentified compute. Within every rung we consider two provers: a **benign** (non-adaptive) prover who runs its workload without regard to the meter, and an **adaptive** (hiding) prover who knows the detector and reshapes its schedule to evade it. Most of what is interesting lies in the gap between them.

Governance reading. A passive meter certifies efficient iteration-structured training against cheap, fixed tests only within a bounded de-periodicisation budget; a prover free to vary real work escapes those tests for nothing. Defeating the order tracker on top of that requires either a systems and learning cost we could not measure or a degradation of the meter channel itself (Section 5.2). The results of this section are therefore necessary conditions and an upper bound on hideability rather than a proof of training: they bound what a prover can hide from a fixed detector cheaply. The complementary direction — a lower bound on what hiding must cost, which no detection sweep can supply — is available for one attack family against a fixed verifier (Section 7.1) under the cost-of-hiding assumption, and the distortion it requires can be priced against the same anchors used here. The tracking class dominates the fixed class at every cell we measured, so that threshold is a valid — but loose — reading of what a tracker would require: it is read off the weaker detector, so it understates the true requirement. Establishing it against the tracking class directly, rather than inheriting it from the measured dominance, remains open.

Table 5. The claim ladder. Each rung answers a strictly more governance-relevant question at a stated cost in verifier information; each rung’s hard ceiling motivates the next. Within each rung we consider a benign (non-adaptive) and an adaptive (hiding) prover. Table 6 closes each rung out against the evidence.

Rung	Claim	Additional requirement	Hard ceiling
1. Structural evidence	Order-coherent power variation exists under a specified structural null	External time-resolved trace only	Periodic physical confounds; structural-null misspecification; strong superposition or dilution; active power shaping; cadence moved outside the searched family; die-to-meter attenuation
2. Conditional classification	The trace is more training-like than the stated inference population	A representative or explicitly modelled inference null	Domain shift; semantic non-identifiability (training-shaped execution, benign or deliberate, without training)
3. Active physical authentication	A load behind the meter responded to an unpredictable verifier challenge	Prover participation; a challenge-bound schedule	A challenge-aware dummy load; no semantic binding by power alone
4. Work verification	The physical response is tied to declared training state transitions	A sound, fresh transcript/proof and an explicit binding assumption	The remote-proof / co-location gap, unless the assumed primitive closes it

Governance reading. What a regulator can buy today is graded. Rung 1 buys physical evidence of an efficient iteration-structured compute schedule from an external meter alone, robust to a wander-and-drift adversary but not to a prover who varies real work or degrades the meter. Rung 2 buys the further claim that the schedule is more training-like than a stated inference baseline — but not that the computation is training rather than a schedule-matched mimic, benign or deliberate, which is a limit of the passive channel, not of the method. That limit runs both ways: it caps what a positive detection proves as well as what a counterfeit could fake. Separating training from a mimic, and binding a physical response to declared work, both require moving off the passive channel to the active protocol of Appendix H and the cryptographic primitive it assumes. The passive meter is thus a cheap, deployable front line that narrows the question and prices the cheap evasions, not a stand-alone proof of training.

The problem of

The problem of X is well established

Inferring which activities are running from one aggregate meter is the founding problem of non-intrusive load monitoring (Hart, 1992; Zoha et al., 2012), with a power-side-channel malware detector as its closest security-flavoured descendant (Clark et al., 2013). Deciding whether a source is active against a transmitter trying to stay below the detection floor is spectrum sensing, whose taxonomy of energy, matched-filter and cyclostationary detectors structures our own comparison (Yücek & Arslan, 2009).

Tracking a narrowbandline whose frequency wanders is passive sonar’s lofargram problem (Streit & Barrett, 1990), carried into continuous-gravitational-wave searches (Suvorova et al., 2016; Bayley et al., 2019), and undoing that wander before testing is tachless order tracking (Bonnardot et al., 2005).

Each of these literatures assumes something the governance setting removes:

load monitoring assumes indifferent appliances with clean switching edges;

side-channel analysis, from differential power analysis onward (Kocher et al., 1999; Chari et al., 2002), assumes a proximate high-bandwidth probe and a cooperative target;

the line trackers assume a stationary noise floor with no adversary behind it.

None of them supplies a calibrated test for a quasi-periodic, frequency-wandering signature observed through a low-rate filtered meter, and none prices what it costs the source to remove it.

The adversarial half does have an owner. Pricing concealment against detectability is the warden-versus-covert-transmitter game of covert communication (Bash et al., 2013) and the relative-entropy security measure of information-theoretic steganography (Cachin, 2004): both bound how detectable a hidden activity remains as a function of how constrained it

Table 6. The claim ladder closed out: what this paper establishes at each rung, where, and the hard ceiling it hits. Mirrors Table 5.

Rung	What the paper establishes	Evidence	Demonstrated ceiling
1. Structural evidence	Order-tracked cyclic structure is detected against stationary and controller-like nulls; the tracker is Neyman–Pearson optimal against the inference null	Section 5, Appendix C	Fast cadence wander (>0.8 Hz on 300 s); attributing structure to training needs the nuisance model; the meter must meet Section 5.2
2. Conditional classification	The trace is more training-like than the stated inference null (area under the curve 1.00, stable across domain shift)	Section F.2	Semantic non-identifiability: non-training loads with a training-shaped schedule score as training, benign ones included; Section 7.1 shows no passive detector can do better
Frontier	Fixed tests collapse under $\approx$ zero-cost variable real work while the tracker holds; costly de-periodicisations buy no concealment	Section 6	The tracker’s far end, at a systems and learning cost we could not measure
Cost of hiding	Hiding to detectability ε requires a distortion $\sigma^*(\varepsilon)$ that the measured anchors price at tens of percent of throughput, rising above 100% (i.i.d. period jitter)	Section 7.1, Appendix G	Read off a fixed-bin verifier, so valid against any class containing it, but loose; one attack family; internal to the stated generator; the throughput-or-learning disjunction is an assumption, not a result
3–4. Active authentication	A challenge-synchronous protocol and an assumed binding primitive, specified but not run	Appendix H	Challenge-aware dummy load; the co-location gap

is, and both do so for *every* verifier at once rather than for one detector. Their results assume a transmitter free to shape its signal. A training schedule is not free: its communication phase has a physical duration that does not shrink when the compute phase does, and each distortion of the cadence costs throughput or learning efficiency. That constraint is what makes a cost-of-hiding bound possible here, and no existing bound covers it.

What is missing is a method that turns a low-rate external power trace into a calibrated test for iteration-structured training, prices its evasion against an adversary who knows the test, and states what the channel cannot certify at any price. Without one, “power monitoring” enters policy proposals as a binary capability: no one can write the meter specification a regulator would have to procure, budget the evasions a prover would actually attempt, or say which governance claims a positive detection does and does not support.

and U.S. Executive Order 14110 briefly operationalised one, requiring reports of runs above 10^{26} FLOP (The White House, 2023).

Every such scheme presupposes that a regulator can tell when training is happening on monitored hardware. Three channels could tell it. Prover-owned telemetry is informative but conditional on tamper-resistant read paths that do not yet exist (Rahman & Tajdari, 2026; O’Gara et al., 2025); challenge–response protocols remove that trust anchor but require the prover’s participation (Monfared et al., 2026); and a verifier-owned probe on an accelerator’s supply conductors requires neither, at the price of physical access to the board and continuous on-site presence (Gargiulo & Kulp, 2026).

Multi GPU synchronous AI training of the type that typifies frontier AI models alternates a compute phase, a forward backward pass with high power requirements, with a communication gradient all reduce phase where the GPUs all sync and power requirements drop. [what about async training? maybe we need to mention this somewhere? perhaps in limitations?]

In practice the problem is harder as the probe is at a higher level i.e. more remote a meter three to four orders of magnitude slower than the probe, behind a power-delivery path that low-passes and dilutes the signal, reading a prover who knows it is being read.